

LFD259

Kubernetes for Developers

Version 2022-11-08



Version 2022-11-08

© Copyright the Linux Foundation 2022. All rights reserved.

© Copyright the Linux Foundation 2022. All rights reserved.

The training materials provided or developed by The Linux Foundation in connection with the training services are protected by copyright and other intellectual property rights.

Open source code incorporated herein may have other copyright holders and is used pursuant to the applicable open source license.

The training materials are provided for individual use by participants in the form in which they are provided. They may not be copied, modified, distributed to non-participants or used to provide training to others without the prior written consent of The Linux Foundation.

No part of this publication may be reproduced, photocopied, stored on a retrieval system, or transmitted without express prior written consent.

Published by:

the **Linux Foundation**

<https://www.linuxfoundation.org>

No representations or warranties are made with respect to the contents or use of this material, and any express or implied warranties of merchantability or fitness for any particular purpose or specifically disclaimed.

Although third-party application software packages may be referenced herein, this is for demonstration purposes only and shall not constitute an endorsement of any of these software applications.

Linux is a registered trademark of Linus Torvalds. Other trademarks within this course material are the property of their respective owners.

If there are any questions about proper and fair use of the material herein, please go to <https://trainingsupport.linuxfoundation.org>.

Nondisclosure of Confidential Information

“Confidential Information” shall not include any of the following, even if marked confidential or proprietary: (a) information that relates to the code base of any open source or open standards project (collectively, “Open Project”), including any existing or future contribution thereto; (b) information generally relating or pertaining to the formation or operation of any Open Project; or (c) information relating to general business matters involving any Open Project.

This course does not include confidential information, nor should any confidential information be divulged in class.

Contents

2	Kubernetes Architecture	1
3	Build	17
4	Design	31
5	Deployment Configuration	47
6	Understanding Security	67
7	Exposing Applications	85
8	Application Troubleshooting	103

Chapter 2

Kubernetes Architecture



Exercise 2.1: Overview and Preliminaries

We will create a two-node **Ubuntu 20.04** cluster. Using two nodes allows an understanding of some issues and configurations found in a production environment. Currently 2 vCPU and 8G of memory allows for quick labs. As we will be working with images it is suggested that you allocate at least **20G** of disk space for the control plane (cp) node, more is better. Other Linux distributions should work in a very similar manner, but have not been tested.

Summary:

- 2 CPU
- 8G memory
- 20G+ disk on control plane node
- Don't use 192.168 network for nodes
- No firewall
- Disable swap
- Disable SELinux and AppArmor



Very Important

Regardless of the platform used (**VirtualBox**, **VMWare**, **AWS**, **GCE** or even bare metal) please remember that security software like **SELinux**, **AppArmor**, and firewall configurations can prevent the labs from working. While not something to do in production consider disabling the firewall and security software.

GCE requires a new VPC to be created and a rule allowing all traffic to be included. The use of **wireshark** can be a helpful place to start with troubleshooting network and connectivity issues if you're unable to open all ports.

The **kubeadm** utility currently requires that swap be turned off on every node. The **swapoff -a** command will do this until the next reboot, with various methods to disable swap persistently. Cloud providers typically deploy instances with swap disabled.

Download shell scripts and YAML files

To assist with setting up your cluster please download the tarball of shell scripts and YAML files. The `k8scp.sh` and `k8sWorker.sh` scripts deploy a Kubernetes cluster using **kubeadm** and use Project Calico for networking. Should the file not be found you can always use a browser to investigate the parent directory.

```
$ wget https://training.linuxfoundation.org/cm/LFD259/LFD259_V2022-11-08_SOLUTIONS.tar.xz \
--user=LFtraining --password=Penguin2014
```

```
$ tar -xvf LFD259_V2022-11-08_SOLUTIONS.tar.xz
```

(Note: depending on your software, if you are cutting and pasting the above instructions, the underscores may disappear and be replaced by spaces, so you may have to edit the command line by hand!)

✍ Exercise 2.2: Deploy a New Cluster

Deploy a Control Plane Node using Kubeadm

1. Log into your nodes using **PuTTY** or using **SSH** from a terminal window. Unless the instructor tells you otherwise the user name to use will be **student**. You may need to change the permissions on the `pem` (or `ppk` on windows) file as shown in the following commands. Your file name and node IP address will probably be different.

If using PuTTY, search for instructions on using a `ppk` key to access an instance. Use the **student** username when asked by PuTTY.

```
localTerm:~$ chmod 400 LF-Class.pem
localTerm:~$ ssh -i LF-Class.pem student@WW.XX.YY.ZZ
```

```
student@cp:~$
```

2. Use the **wget** and **tar** commands shown above to download and extract the course tarball to your `cp` node. Install **wget** if the command is not found on your instance.
3. Be sure to type commands found in the exercises, instead of trying to copy and paste. Typing will cause you to learn the commands, whereas copy and paste tends not to be absorbed. Also, use the YAML files included in the course tarball, as whitespace issues can slow down the labs and be discouraging.
4. Review the script to install and begin the configuration of the `cp` kubernetes server. You may need to change the **find** command search directory for your home directory depending on how and where you downloaded the tarball.

A **find** command is shown if you want to locate and copy to the current directory instead of creating the file. Mark the command for reference as it may not be shown for future commands.

```
student@cp:~$ find $HOME -name <YAML File>
student@cp:~$ cp LFD259/<Some Path>/<YAML File> .
```

```
student@cp:~$ find $HOME -name k8scp.sh
```

```
student@cp:~$ more LFD259/SOLUTIONS/s_02/k8scp.sh
```

```
....
```

```
# Install the Kubernetes software, and lock the version
sudo apt update
sudo apt -y install kubelet=1.25.1-00 kubeadm=1.25.1-00 kubectl=1.25.1-00
sudo apt-mark hold kubelet kubeadm kubectl
```

```
# Ensure Kubelet is running
```

```
sudo systemctl enable --now kubelet

# Disable swap just in case
sudo swapoff -a
....
```

- Run the script as an argument to the **bash** shell. You will need the `kubeadm join` command shown near the end of the output when you add the worker/minion node in a future step. Use the **tee** command to save the output of the script, in case you cannot scroll back to find the `kubeadm join` in the script output. Please note the following is one command and then its output.

Using **Ubuntu 20.04** you may be asked questions during the installation. Allow restarts (yes) and use the local, installed software if asked during the update, usually (option 2).

Copy files to your home directory first.

```
student@cp:~$ cp LFD259/SOLUTIONS/s_02/k8scp.sh .
```

```
student@cp:~$ bash k8scp.sh | tee $HOME/cp.out
```

```
<output_omitted>

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

    mkdir -p $HOME/.kube
    sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
    sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

    export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
    https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 10.128.0.93:6443 --token zu7icq.2o0ouw7s044yi9w9 \
    --discovery-token-ca-cert-hash
    ↪ sha256:a26f2166ad39b6aba0fc0dd663ac2b3455af19526bd9fe80af5439255dd
<output_omitted>

NAME          STATUS    ROLES    AGE     VERSION
cp            NotReady control-plane 20s     v1.25.1

Script finished. Move to the next step
```

Deploy a Worker Node

- Open a separate terminal into your **second node**, which will be your worker. Having both terminal sessions allows you to monitor the status of the cluster while adding the second node. Change the color or other characteristic of the second terminal to make it visually distinct from the first. This will keep you from running commands on the incorrect instance, which probably won't work.

Use the previous **wget** command download the tarball to the worker node. Extract the files with **tar** as before. Find and copy the `k8sWorker.sh` file to student's home directory then view it. You should see the same early steps as found in the cp setup script.

```
student@worker:~$ more k8sWorker.sh
```

```
....
# Install the Kubernetes software, and lock the version
sudo apt update
sudo apt -y install kubelet=1.25.1-00 kubeadm=1.25.1-00 kubectl=1.25.1-00
sudo apt-mark hold kubelet kubeadm kubectl

# Ensure Kubelet is running
sudo systemctl enable --now kubelet

# Disable swap just in case
sudo swapoff -a
....
```

7. Run the script on the **second node**. Again please note you may have questions during the update. Allow daemons to restart, type yes, and use the local installed version, usually option 2. For troubleshooting you may want to write output to a file using the **tee** command.

```
student@worker:~$ bash k8sWorker.sh | tee worker.out
```

```
<output_omitted>
```

8. When the script is done the worker node is ready to join the cluster. The `kubeadm join` statement can be found near the end of the `kubeadm init` output on the cp node. It should also be in the file `cp.out` as well. Your nodes will use a different IP address and hashes than the example below. You'll need to pre-pend **sudo** to run the script copied from the cp node. Also note that some non-Linux operating systems and tools insert extra characters when multi-line samples are copied and pasted. Copying one line at a time solves this issue. A helpful command to find a line in a file may be **grep -A2 join cp.out** on the control plane node.

```
student@worker:~$ sudo kubeadm join --token 118c3e.83b49999dc5dc034 \
10.128.0.3:6443 --discovery-token-ca-cert-hash \
sha256:40aa946e3f53e38271bae24723866f56c86d77efb49aede8a70cc189bfe2e1d
```

```
<output_omitted>
```

Configure the Control Plane Node

9. Return to the cp node. Install a text editor. While the lab uses **vim**, any text editor such as **emacs** or **nano** will work. Be aware that Windows editors may have issues with special characters. Also install the **bash-completion** package, if not already installed. Use the locally installed version of a package if asked.

```
student@cp:~$ sudo apt-get install bash-completion vim -y
```

```
<output_omitted>
```

10. We will configure command line completion and verify both nodes have been added to the cluster. The first command will configure completion in the current shell. The second command will ensure future shells have completion. You may need to exit the shell and log back in for command completion to work without error.

```
student@cp:~$ source <(kubectl completion bash)
```

```
student@cp:~$ echo "source <(kubectl completion bash)" >> $HOME/.bashrc
```

11. Verify that both nodes are part of the cluster. And show a Ready state.

```
student@cp:~$ kubectl get node
```

NAME	STATUS	ROLES	AGE	VERSION
cp	Ready	control-plane	4m11s	v1.25.1
worker	Ready	<none>	61s	v1.25.1

12. We will use the **kubectl** command for the majority of work with Kubernetes. Review the help output to become familiar with commands options and arguments.

```
student@cp:~$ kubectl --help
```

```
kubectl controls the Kubernetes cluster manager.

Find more information at:
  https://kubernetes.io/docs/reference/kubectl/overview/

Basic Commands (Beginner):
  create      Create a resource from a file or from stdin.
  expose      Take a replication controller, service,
  deployment or pod and expose it as a new Kubernetes Service
  run         Run a particular image on the cluster
  set         Set specific features on objects

Basic Commands (Intermediate):
<output_omitted>
```

13. With more than 40 arguments, you can explore each also using the `--help` option. Take a closer look at a few, starting with `taint` for example.

```
student@cp:~$ kubectl taint --help
```

```
Update the taints on one or more nodes.

* A taint consists of a key, value, and effect. As an argument
  here, it is expressed as key=value:effect.
* The key must begin with a letter or number, and may contain
  letters, numbers, hyphens, dots, and underscores, up to
  253 characters.
* Optionally, the key can begin with a DNS subdomain prefix
  and a single '/',
  like example.com/my-app
<output_omitted>
```

14. By default the `cp` node will not allow general containers to be deployed for security reasons. This is via a taint. Only containers which tolerate this taint will be scheduled on this node. As we only have two nodes in our cluster we will remove the taint, allowing containers to be deployed on both nodes. This is not typically done in a production environment for security and resource contention reasons. The following command will remove the taint from all nodes, so you should see one success and one not found error. The worker/minion node does not have the taint to begin with. Note the **minus sign** at the end of the command, which removes the preceding value.

```
student@cp:~$ kubectl describe nodes | grep -i taint
```

```
Taints:          node-role.kubernetes.io/control-plane:NoSchedule
Taints:          <node>
```

```
student@cp:~$ kubectl taint nodes --all node-role.kubernetes.io/control-plane-
```

```
node/cp untainted
error: taint "node-role.kubernetes.io/control-plane" not found
```

15. As we transition away from the use of the term "master", you may find a taint still exists. Remove it as well.

```
student@cp:~$ kubectl describe nodes | grep -i taint
```

```
Taints:          node-role.kubernetes.io/master:NoSchedule
Taints:          <node>
```

```
student@cp:~$ kubectl taint nodes --all node-role.kubernetes.io/master-
```

```
node/cp untainted
error: taint "node-role.kubernetes.io/master" not found
```

16. Check that both nodes are without a Taint. If they both are without taint the nodes should now show as Ready. It may take a minute or two for all infrastructure pods to enter Ready state, such that the nodes will show a Ready state.

```
student@cp:~$ kubectl describe nodes | grep -i taint
```

```
Taints:          <none>
Taints:          <none>
```

```
student@cp:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
cp	Ready	control-plane	6m1s	v1.25.1
worker	Ready	<none>	5m31s	v1.25.1

✍ Exercise 2.3: Create a Basic Pod

1. The smallest unit we directly control with Kubernetes is the pod. We will create a pod by creating a minimal YAML file. First we will get a list of current API objects and their APIGROUP. If value is not shown it may not exist, as with SHORTNAMES. Note that pods does not declare an APIGROUP. At the moment this indicates it is part of the stable v1 group.

```
student@cp:~$ kubectl api-resources
```

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
bindings		v1	true	Binding
componentstatuses	cs	v1	false	ComponentStatus
configmaps	cm	v1	true	ConfigMap
endpoints	ep	v1	true	Endpoints
.....				
pods	po	v1	true	Pod
....				

2. From the output we see most are v1 which is used to denote a stable object. With that information we will add the other three required sections for pods such as metadata, with a name, and spec which declares which container image to use and a name for the container. We will create an eight line YAML file. White space and indentation matters. Don't use **Tab**s. There is a `basic.yaml` file available in the tarball, as well as `basic-later.yaml` which shows what the file will become and can be helpful for figuring out indentation.

YAML

basic.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5 spec:
6   containers:
7   - name: webcont
8     image: nginx
```

3. Create the new pod using the recently created YAML file.

```
student@cp:~$ kubectl create -f basic.yaml
```

```
pod/basicpod created
```

- Make sure the pod has been created then use the **describe** sub-command to view the details. Among other values in the output you should be about to find the image and the container name.

```
student@cp:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
basicpod	1/1	Running	0	23s

```
student@cp:~$ kubectl describe pod basicpod
```

```
Name:          basicpod
Namespace:     default
Priority:       0
<output_omitted>
```

- Shut down the pod and verify it is no longer running.

```
student@cp:~$ kubectl delete pod basicpod
```

```
pod "basicpod" deleted
```

```
student@cp:~$ kubectl get pod
```

```
No resources found in default namespace.
```

- We will now configure the pod to expose port 80. This configuration does not interact with the container to determine what port to open. We have to know what port the process inside the container is using, in this case port 80 as a web server. Add two lines to the end of the file. Line up the indentation with the `image` declaration.

```
student@cp:~$ vim basic.yaml
```

YAML

basic.yaml

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: basicpod
5  spec:
6    containers:
7      - name: webcont
8        image: nginx
9        ports:
10         - containerPort: 80
```

#<--Add this and following line

- Create the pod and verify it is running. Use the `-o wide` option to see the internal IP assigned to the pod, as well as `NOMINATED NODE`, which is used by the scheduler and `READINESS GATES` which show if experimental features are enabled. Using **curl** and the pods IP address you should get the default nginx welcome web page.

```
student@cp:~$ kubectl create -f basic.yaml
```

```
pod/basicpod created
```

```
student@cp:~$ kubectl get pod -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE READINESS GATES						
basicpod	1/1	Running	0	9s	192.168.1.3	cp
<none>		<none>				

```
student@cp:~$ curl http://192.168.1.3
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>

<output_omitted>
```

```
student@cp:~$ kubectl delete pod basicpod
```

```
pod "basicpod" deleted
```

8. We will now create a simple service to expose the pod to other nodes and pods in the cluster. The service YAML will have the same four sections as a pod, but different spec configuration and the addition of a selector. Again, copy of the example from the tarball instead of typing the file by hand.

```
student@cp:~$ vim basicservice.yaml
```

YAML

basicservice.yaml

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: basicservice
5 spec:
6   selector:
7     type: webserver
8   ports:
9     - protocol: TCP
10      port: 80
```

9. We will also add a label to the pod and a selector to the service so it knows which object to communicate with.

```
student@cp:~$ vim basic.yaml
```

YAML

basic.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5   labels:                                #<-- Add this line
6     type: webserver                      #<-- and this line which matches selector
7 spec:
8   ....
```

10. Create the new pod and service. Verify both have been created. We will learn details of the output in later chapters.

```
student@cp:~$ kubectl create -f basic.yaml
```

```
pod/basicpod created
```

```
student@cp:~$ kubectl create -f basicservice.yaml
```

```
service/basicservice created
```

```
student@cp:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
basicpod	1/1	Running	0	110s

```
student@cp:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
basicservice	ClusterIP	10.96.112.50	<none>	80/TCP	14s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4h

11. Test access to the web server using the CLUSTER-IP for the basicservice.

```
student@cp:~$ curl http://10.96.112.50
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>

<output_omitted>
```

12. We will now expose the service to outside the cluster as well. Delete the service, edit the file and add a type declaration.

```
student@cp:~$ kubectl delete svc basicservice
```

```
service "basicservice" deleted
```

```
student@cp:~$ vim basicservice.yaml
```

YAML

basicservice.yaml

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: basicservice
5 spec:
6   selector:
7     type: webserver
8   type: NodePort      #<--Add this line
9   ports:
10  - protocol: TCP
11    port: 80
```

13. Create the service again. Note there is a different TYPE and CLUSTER-IP and also a high-numbered port.

```
student@cp:~$ kubectl create -f basicservice.yaml
```

```
service/basicservice created
```

```
student@cp:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
basicservice	NodePort	10.100.139.155	<none>	80:31514/TCP	3s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	47h

14. Using the public IP address of the node and the high port you should be able to test access to the webserver. In the example below the public IP is 35.238.3.83, as reported by a **curl** to ifconfig.io. Your IP will be different. The high port will also probably be different. Note that testing from within a GCE or AWS node will not work. Use a local to you terminal or web browser to test.

```
student@cp:~$ curl ifconfig.io
```

```
35.238.3.83
```

```
local$ curl http://35.238.3.83:31514
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>
```

✍ Exercise 2.4: Multi-Container Pods

Using a single container per pod allows for the most granularity and decoupling. There are still some reasons to deploy multiple containers, sometimes called *composite containers*, in a single pod. The secondary containers can handle logging or enhance the primary, the *sidecar* concept, or acting as a proxy to the outside, the *ambassador* concept, or modifying data to meet an external format such as an *adapter*. All three concepts are secondary containers to perform a function the primary container does not.

1. We will add a second container to the pod to handle logging. Without going into details of how to use **fluentd** we will add a logging container to the existing pod from its own repository. The second container would act as a *sidecar*. At this state we will just add the second container and verify it is running. In the **Deployment Configuration** chapter we will continue to work on this pod by adding persistent storage and configure **fluentd** via a `configMap`.

Edit the YAML file and add a **fluentd** container. The dash should line up with the previous container dash. At this point a name and image should be enough to start the second container.

```
student@cp:~$ vim basic.yaml
```

YAML

basic.yaml

```
1  ....
2  containers:
3  - name: webcont
4    image: nginx
5    ports:
6  - containerPort: 80
7  - name: fdlogger
8    image: fluent/fluentd
```

2. Delete and create the pod again. The commands can be typed on a single line, separated by a semicolon. This time you should see 2/2 under the READY column. You should also find information on the **fluentd** container inside of the `kubectl describe` output.

```
student@cp:~$ kubectl delete pod basicpod ; kubectl create -f basic.yaml
```

```
pod "basicpod" deleted
pod/basicpod created
```

```
student@cp:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
basicpod	2/2	Running	0	2m8s

```
student@cp:~$ kubectl describe pod basicpod
```

```
Name:          basicpod
Namespace:     default
Priority:       0
Node:          cp/10.128.0.11
...

fdlogger:
  Container ID:  docker://f0649457217f00175ce9aec35022d0b238b9b...
  Image:         fluent/fluentd
...
```

3. For now shut down the pod. We will use it again in a future exercise.

```
student@cp:~$ kubectl delete pod basicpod
```

```
pod "basicpod" deleted
```

Exercise 2.5: Create a Simple Deployment

Creating a pod does not take advantage of orchestration abilities of Kubernetes. We will now create a Deployment which gives us scalability, reliability, and updates.

1. Now run a containerized webserver **nginx**. Use **kubectl create** to create a simple, single replica deployment running the nginx web server. It will create a single pod as we did previously but with new controllers to ensure it runs as well as other features.

```
student@cp:~$ kubectl create deployment firstpod --image=nginx
```

```
deployment.apps/firstpod created
```

2. Verify the new deployment exists and the desired number of pods matches the current number. Using a comma, you can request two resource types at once. The **Tab** key can be helpful. Type enough of the word to be unique and press the **Tab** key, it should complete the word. The deployment should show a number 1 for each value, such that the desired number of pods matches the up-to-date and running number. The pod should show zero restarts.

```
student@cp:~$ kubectl get deployment,pod
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/firstpod	1/1	1	1	2m42s

NAME	READY	STATUS	RESTARTS	AGE
pod/firstpod-7d88d7b6cf-lrsbk	1/1	Running	0	2m42s

3. View the details of the deployment, then the pod. Work through the output slowly. Knowing what a healthy deployment and looks like can be helpful when troubleshooting issues. Again the **Tab** key can be helpful when using long auto-generated object names. You should be able to type `firstpod` and the name will complete when viewing the pod.

```
student@cp:~$ kubectl describe deployment firstpod
```

```
Name:                firstpod
Namespace:           default
CreationTimestamp:    Tue, 21 Sep 2021 16:36:29 +0000
Labels:              app=firstpod
Annotations:         deployment.kubernetes.io/revision=1
Selector:            app=firstpod
Replicas:            1 desired | 1 updated | 1 total | 1 available....
StrategyType:        RollingUpdate
MinReadySeconds:     0
<output_omitted>
```

```
student@cp:~$ kubectl describe pod firstpod-6bb4574d94-rqk76
```

```
Name:                firstpod-6bb4574d94-rqk76
Namespace:           default
Priority:             0
PriorityClassName:    <none>
Node:                cp/10.128.0.2
Start Time:          Tue, 21 Sep 2021 16:36:29 +0000
Labels:              pod-template-hash=2660130850
                    app=firstpod
Annotations:         cni.projectcalico.org/podIP: 192.168.200.65/32
Status:              Running
IP:                  192.168.200.65
Controlled By:       ReplicaSet/firstpod-6bb4574d94
<output_omitted>
```

4. Note that the resources are in the default namespace. Get a list of available namespaces.

```
student@cp:~$ kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	20m
kube-node-lease	Active	20m
kube-public	Active	20m
kube-system	Active	20m

5. There are four default namespaces. Look at the pods in the `kube-system` namespace.

```
student@cp:~$ kubectl get pod -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-node-5ftrr	2/2	Running	0	24m
calico-node-f7zrw	2/2	Running	0	21m
coredns-fb8b8dccf-cmkds	1/1	Running	0	24m
coredns-fb8b8dccf-grltk	1/1	Running	0	24m
etcd-v141-r24p	1/1	Running	0	23m

<output_omitted>

6. Now look at the pods in a namespace that does not exist. Note you do not receive an error.

```
student@cp:~$ kubectl get pod -n fakenamespace
```



```
No resources found in fakenamespaces namespace.
```

7. You can also view resources in all namespaces at once. Use the `--all-namespaces` options to select objects in all namespaces at once.

```
student@cp:~$ kubectl get pod --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	firstpod-69cfd8d9-kj6ql	1/1	Running	0	44m
kube-system	calico-node-5ftrr	2/2	Running	0	92m
kube-system	calico-node-f7zrw	2/2	Running	0	89m
kube-system	coredns-fb8b8dcf-cmkds	1/1	Running	0	92m
<output_omitted>					

8. View several resources at once. Note that most resources have a short name such as `rs` for ReplicaSet, `po` for Pod, `svc` for Service, and `ep` for endpoint. Note the endpoint still exists after we deleted the pod.

```
student@cp:~$ kubectl get deploy,rs,po,svc,ep
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/firstpod	1/1	1	1	4m

NAME	DESIRED	CURRENT	READY...
replicaset.apps/firstpod-6bb4574d94-rqk76	1	1	1

NAME	READY	STATUS	RESTARTS	AGE
pod/firstpod-6bb4574d94-rqk76	1/1	Running	0	4m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/basicservice	NodePort	10.108.147.76	<none>	80:31601/TCP	21m
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	21m

NAME	ENDPOINTS	AGE
endpoints/basicservice	<none>	21m
endpoints/kubernetes	10.128.0.3:6443	21m

9. Delete the ReplicaSet and view the resources again. Note that the age on the ReplicaSet and the pod it controls is now less than a minute of age. The deployment operator started a new ReplicaSet operator when we deleted the existing one. The new ReplicaSet started another pod when the desired spec did not match the current status.

```
student@cp:~$ kubectl delete rs firstpod-6bb4574d94-rqk76
```

```
replicaset.apps "firstpod-6bb4574d94-rqk76" deleted
```

```
student@cp:~$ kubectl get deployment,rs,po,svc,ep
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/firstpod	1/1	1	1	7m

NAME	DESIRED	CURRENT...
replicaset.apps/firstpod-6bb4574d94-rqk76	1	1

NAME	READY	STATUS	RESTARTS	AGE
pod/firstpod-7d99ffc75-p9hbw	1/1	Running	0	12s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	24m

NAME	ENDPOINTS	AGE
endpoints/kubernetes	10.128.0.2:6443	80m

```
endpoints/basicservice <none> 21m
```

10. This time delete the top-level controller. After about 30 seconds for everything to shut down you should only see the cluster service and endpoint remain for the cluster and the service we created.

```
student@cp:~$ kubectl delete deployment firstpod
```

```
deployment.apps "firstpod" deleted
```

```
student@cp:~$ kubectl get deployment,rs,po,svc,ep
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/basicservice	NodePort	10.108.147.76	<none>	80:31601/TCP	35m
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	24m

NAME	ENDPOINTS	AGE
endpoints/basicservice	<none>	21m
kubernetes	10.128.0.3:6443	24m

11. As we won't need it for a while, delete the basicservice service as well.

```
student@cp:~$ kubectl delete svc basicservice
```

```
service "basicservice" deleted
```

Exercise 2.6: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

1. Using a browser go to <https://www.cncf.io/certification/ckad/> and read through the program description.
2. In the **Exam Resources** section open the [Curriculum Overview](#) and [Candidate Handbook](#) in new tabs. Both of these should be read and understood prior to sitting for the exam.
3. Navigate to the [Curriculum Overview](#) tab. You should see links for domain information for various versions of the exam. Select the latest version, such as **CKAD_Curriculum_V1.23.1.pdf**. The versions you see may be different. You should see a new page showing a PDF.
4. Read through the document. Be aware that the term Understand, such as Understand Services, is more than just knowing they exist. In this case expect it to also mean create, update, and troubleshoot.
5. Locate the **Application Design and Build** section. If you review the lab, you will see we have covered some of these steps such as multi-container Pod design. Again, please note this document will change, distinct from this book. **It remains your responsibility to check for changes in the online document.** They may change on an irregular basis.

Certified Kubernetes Application Developer (CKAD) Exam Curriculum

This document provides the curriculum outline of the Knowledge, Skills and Abilities that a Certified Kubernetes Application Developer (CKAD) can be expected to demonstrate.

CKAD Curriculum

20% - Application Design and Build

- Define, build and modify container images
- Understand Jobs and CronJobs
- Understand multi-container Pod design patterns (e.g. sidecar, init and others)
- Utilize persistent and ephemeral volumes

25% - Application Environment, Configuration and Security

- Discover and use resources that extend Kubernetes (CRD)
- Understand authentication, authorization and admission control
- Understanding and defining resource

Figure 2.1: Application Design and Build Domain

6. Navigate to the **Candidate Handbook** tab. You are strongly encourage to read and understand this entire document prior to taking the exam. Again, please note this document will change, distinct from this book. **It remains your responsibility to check for changes in the online document.** They may change on an irregular and unannounced basis.
7. Find the **Important Instructions: CKA and CKAD** section in the document. Read through the section, locate current Kubernetes version for exams, such as v1.22, and find the **Resources allowed during exam** section.
8. Note the domains and subdomains you can use during the exam, with some stated conditions.

Resources allowed during exam

During the exam, candidates may:

- review the Exam content instructions that are presented in the command line terminal
- review Documents installed by the distribution (i.e. /usr/share and its subdirectories)
- use their Chrome or Chromium browser to open one additional tab in order to access assets at: <https://kubernetes.io/docs/>, <https://github.com/kubernetes/>, <https://kubernetes.io/blog/> and their subdomains. This includes all available language translations of these pages (e.g. <https://kubernetes.io/zh/docs/>)

No other tabs may be opened and no other sites may be navigated to (including <https://discuss.kubernetes.io/>).

The allowed sites above may contain links that point to external sites. It is the responsibility of the candidate not to click on any links that cause them to navigate to a domain that is not allowed.

Exam Details

System Requirements to take...

Acceptable Testing Location

ID Requirements to take the ...

Sanctioned Countries

Resources allowed during ex...

Exam Technical Instructions

CKA & CKAD Environment

Additional Resources

Figure 2.2: Exam Handbook Resources Allowed

9. Using only the allowed browser, URLs, and subdomains search for and bookmark a YAML example to create and configure a basic pod. Ensure it works for the version of the exam you are taking. URLs may change, plan on checking each book mark prior to taking the exam. Practice finding the good, working YAML quickly.

10. Using a timer and bookmarked YAML files see how long it takes you to create and verify. Try it again and see how much faster you can complete and test each step:
 - A new pod with the **nginx** image. Showing all containers running and a Ready status.
 - A new service exposing the pod as a **nodePort**, which presents a working webserver configured in the previous step.
 - Update the pod to run the **nginx:1.11-alpine** image and re-verify you can view the webserver via a nodePort.
11. Find and use the `architecture-review1.yaml` file included in the course tarball. Your path, such as course number, may be different than the one in the example below. Use the **find** output. Determine if the pod is running. Fix any errors you may encounter. The use of **kubectl describe** may be helpful.

```
student@cp:~$ find $HOME -name architecture-review1.yaml
```

```
<some_long_path>/architecture-review1.yaml
```

```
student@cp:~$ cp <copy-paste-from-above> .
```

```
student@cp:~$ kubectl create -f architecture-review1.yaml
```

12. Remove any pods or services you may have created as part of the review before moving on to the next section. For example:

```
student@cp:~$ kubectl delete -f architecture-review1.yaml
```

Chapter 3

Build



Exercise 3.1: Deploy a New Application

Overview

In this lab we will deploy a very simple **Python** application, test it using **Podman**, ingest it into Kubernetes and configure probes to ensure it continues to run. This lab requires the completion of the previous lab, the installation and configuration of a Kubernetes cluster.

Note that **Podman** and **crictl** use the same syntax as **docker**, so the labs should work the same way if you happen to be using Docker instead.

Working with A Simple Python Script

1. Install python on your cp node. It may already be installed, as is shown in the output below.

```
student@cp:~$ sudo apt-get -y install python3
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
python3 is already the newest version (3.8.2-0ubuntu2).
The following package was automatically installed and is no longer required:
  libnuma1
Use 'sudo apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.
```

2. Locate the python binary on your system.

```
student@cp:~$ which python3
```

```
/usr/bin/python3
```

3. Create and change into a new directory. The Podman build process pulls everything from the current directory into the image file by default. Make sure the chosen directory is empty.

```
student@cp:~$ mkdir app1
student@cp:~$ cd app1
student@cp:~/app1$ ls -l
```

```
total 0
```

4. Create a simple python script which prints the time and hostname every 5 seconds. There are six commented parts to this script, which should explain what each part is meant to do. The script is included with others in the course tar file, you may consider using the **find** command used before to find and copy over the file.

While the command shows **vim** as an example other text editors such as **nano** work just as well.

```
student@cp:~/app1$ vim simple.py
```



simple.py

```
1  #!/usr/bin/python3
2  ## Import the necessary modules
3  import time
4  import socket
5
6  ## Use an ongoing while loop to generate output
7  while True :
8
9  ## Set the hostname and the current date
10     host = socket.gethostname()
11     date = time.strftime("%Y-%m-%d %H:%M:%S")
12
13     ## Convert the date output to a string
14     now = str(date)
15
16     ## Open the file named date in append mode
17     ## Append the output of hostname and time
18     f = open("date.out", "a" )
19     f.write(now + "\n")
20     f.write(host + "\n")
21     f.close()
22
23     ## Sleep for five seconds then continue the loop
24     time.sleep(5)
```

5. Make the file executable and test that it works. Use Ctrl-C to interrupt the while loop after 20 or 30 seconds. The output will be sent to a newly created file in your current directory called `date.out`.

```
student@cp:~/app1$ chmod +x simple.py
student@cp:~/app1$ ./simple.py
```

```
^CTraceback (most recent call last):
  File "./simple.py", line 42, in <module>
    time.sleep(5)
KeyboardInterrupt
```

6. Verify the output has node name and timestamp.

```
student@cp:~/app1$ cat date.out
```

```
2018-03-22 15:51:38
cp
2018-03-22 15:51:43
cp
```

```
2018-03-22 15:51:48
cp
<output_omitted>
```

7. Create a text file named `Dockerfile`.



Very Important

The name is important: it cannot have a suffix.

We will use three statements, `FROM` to declare which version of Python to use, `ADD` to include our script and `CMD` to indicate the action of the container. Should you be including more complex tasks you may need to install extra libraries, shown commented out as `RUN pip install` in the following example.

```
student@cp:~/app1$ vim Dockerfile
```



Dockerfile

```
FROM python:3
ADD simple.py /
## RUN pip install pystrich
CMD [ "python", "./simple.py" ]
```

8. Build the container. The output below shows the end-build as necessary software was downloaded. You will need to use **sudo** in order to run this command. After the four step process completes the last lines of output should indicate success. Note the dot (.) at the end of the command indicates the current directory.

The **podman** command has been built to replace all of the functionality of **docker**, and should accept the same syntax. As with any open source, fast changing project there could be slight differences. You may note the process takes almost a minute to finish, with a pause or two in output. Some may alias **docker** to **podman**.

Choose to use the `docker.io` version of python.

```
student@cp:~/app1$ sudo podman build -t simpleapp .
```

```
STEP 1/3: FROM python:3
Resolved "python" as an alias (/etc/containers/registries.conf.d/000-shortnames.conf)
Trying to pull docker.io/library/python:3...
Getting image source signatures
Copying blob 0c6b8ff8c37e done
Copying blob 724cfd2dc19b done
Copying blob 808edda3c2e8 done
Copying blob 461bb1d8c517 done
Copying blob e6d3e61f7a50 done
Copying blob 412caad352a3 done
Copying blob 1bb6570cd7ac done
Copying blob aca06d6d45b1 done
Copying blob 678714351737 done
Copying config dfce7257b7 done
Writing manifest to image destination
Storing signatures
STEP 2/3: ADD simple.py /
--> 0e0bd8b3671
STEP 3/3: CMD [ "python", "./simple.py" ]
COMMIT simpleapp
--> 5b825f8d5f3
Successfully tagged localhost/simpleapp:latest
5b825f8d5f3a69660f2278a7119716e2a0fc1a5756d066800429f8030f83e978
```

- Verify you can see the new image among others downloaded during the build process, installed to support the cluster, or you may have already worked with. The newly created `simpleapp` image should be listed first.

```
student@cp:~/app1$ sudo podman images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
localhost/simpleapp	latest	8dd4fa975329	39 seconds ago	925 MB
docker.io/calico/node	v3.19.0	b0744cc52c19	12 days ago	155 MB
docker.io/calico/pod2daemon-flexvol	v3.19.0	a5decf77918d	12 days ago	21.8 MB
docker.io/calico/cni	v3.19.0	3d17cd6307a4	12 days ago	146 MB

<output_omitted>

- Use **sudo podman** to run a container using the new image. While the script is running you won't see any output and the shell will be occupied running the image in the background. After 30 seconds use **ctrl-c** to interrupt. The local `date.out` file will not be updated with new times, instead that output will be a file of the container image.

```
student@cp:~/app1$ sudo podman run localhost/simpleapp
```

```
^CTraceback (most recent call last):
  File "./simple.py", line 24, in <module>
    time.sleep(5)
KeyboardInterrupt
```

- Locate the newly created `date.out` file. The following command should show two files of this name, the one created when we ran `simple.py` and another under `/var/lib/containers` when run via a podman or crio container.

```
student@cp:~/app1$ sudo find / -name date.out
```

```
/home/student/app1/date.out
/var/lib/containers/storage/overlay/
0dea104afa098f608dff06b17b2196d0ba12d09a775243781862abf016e3378a/diff/date.out
```

- View the contents of the `date.out` file created via Podman. Note the need for **sudo** as Podman created the file this time, and the owner is `root`. The long name is shown on several lines in the example, but would be a single line when typed or copied.

```
student@cp:~/app1$ sudo tail \
/var/lib/containers/storage/overlay/
0dea104afa098f608dff06b17b2196d0ba12d09a775243781862abf016e3378a/diff/date.out
```

```
2018-03-22 16:13:46
53e1093e5d39
2018-03-22 16:13:51
53e1093e5d39
2018-03-22 16:13:56
53e1093e5d39
```

✍ Exercise 3.2: Configure A Local Repo

While we could create an account and upload our application to <https://hub.docker.com> or <https://artifacthub.io/>, thus sharing it with the world, we will instead create a local repository and make it available to the nodes of our cluster.

- Create a simple registry using the `easyregistry.yaml` file included in the course tarball. Use the path returned by **find**, which may be different than the one found in the output below.

```
student@cp:~/app1$ find $HOME -name easyregistry.yaml
```



```
<some_long_path>/easyregistry.yaml
```

```
student@cp:~/app1$ kubectl create -f <path_from_output_above>
```

```
service/nginx created
service/registry created
deployment.apps/nginx created
persistentvolumeclaim/nginx-claim0 created
deployment.apps/registry created
persistentvolumeclaim/registry-claim0 created
persistentvolume/vol1 created
persistentvolume/vol2 created
```

2. Take note of the ClusterIP for the new registry service. In the example below it is 10.97.40.62

```
student@cp:~/app1$ kubectl get svc | grep registry
```

```
registry          ClusterIP   10.97.40.62   <none>         5000/TCP,8080/TCP   5m35s
```

3. Verify the repo is working. Please note that if the connection hangs it may be due to a firewall issue. If running your nodes using GCE ensure your instances are using VPC setup and all ports are allowed. If using AWS also make sure all ports are being allowed.

Edit the IP address to that of your localrepo service found in the previous command, and the listed port of 5000

```
student@cp:~/app1$ curl 10.97.40.62:5000/v2/_catalog
```

```
{"repositories": []}
```

4. Configure **podman** to work with non-TLS repos. Edit the `/etc/containers/registries.conf` to uncomment and modify some entries. Ensure your location is within double-quotes as you may get an "Invalid float value" otherwise.

```
student@cp:~/app1$ sudo vim /etc/containers/registries.conf
```

```
....
# # An array of host[:port] registries to try when pulling an unqualified image, in order.
unqualified-search-registries = ["docker.io", "quay.io"]
#
[[registry]]                                #<-- Uncomment
# # The "prefix" field is used to choose the relevant [[registry]] TOML table;
# # (only) the TOML table with the longest match for the input image name
# # (taking into account namespace/repo/tag/digest separators) is used.
....
# # If true, unencrypted HTTP as well as TLS connections with untrusted
# # certificates are allowed.
insecure = true                             #<-- Uncomment/edit to true
....
# # wildcarded format: "*.example.com". In this case, the input reference will
# # be used as-is without any rewrite.
location = "10.97.40.62:5000"               #<-- Uncomment/edit and add IP:PORT
#
# # (Possibly-partial) mirrors for the "prefix"-rooted namespace.
....
```

5. Restart **crio** and make sure it is still running. Should **crio** fail to restart, check for typos and the correct IP address for the location. Sometimes a reboot works to bring crio to typical operation.

```
student@cp:~/app1$ sudo systemctl restart crio
```

```
student@cp:~/app1$ sudo systemctl status crio
```

```
....
Active: active (running) since Wed 2021-05-12 20:48:48 UTC; 2s ago
....
```

6. It may be easier to set the repository as a variable for these labs. Edit the following line to use your service ClusterIP, and port 5000.

```
student@cp:~/app1$ export repo=10.97.40.62:5000
```

```
student@cp:~/app1$ echo "export repo=10.97.40.62:5000" >> $HOME/.bashrc
```

7. Download and tag a typical image from hub.docker.com. Tag the image using the IP and port of the registry, via the \$repo variable.

```
student@cp:~/app1/$ sudo podman pull alpine
```

```
Resolved "alpine" as an alias (/etc/containers/registries.conf.d/000-shortnames.conf)
Trying to pull docker.io/library/alpine:latest...
Getting image source signatures
Copying blob 540db60ca938 done
Copying config 6dbb9cc540 done
Writing manifest to image destination
Storing signatures
6dbb9cc54074106d46d4ccb330f2a40a682d49dda5f4844962b7dce9fe44aaec
```

```
student@cp:~/app1$ sudo podman tag alpine $repo/tagtest
```

8. Push the newly tagged image to your local registry. If you receive an error about an HTTP request to an HTTPS client check that you edited the `/etc/containers/registry.conf` file correctly and restarted the service.

```
student@cp:~/app1$ sudo podman push $repo/tagtest
```

```
Getting image source signatures
Copying blob b2d5eeeaba3a done
Copying config 6dbb9cc540 done
Writing manifest to image destination
Storing signatures
```

9. We will test to make sure we can also pull images from our local repository. Begin by removing the local cached images.

```
student@cp:~/app1$ sudo podman image rm alpine
```

```
Untagged: docker.io/library/alpine:latest
```

```
student@cp:~/app1$ sudo podman image rm $repo/tagtest
```

```
Untagged: 10.97.40.62:5000/tagtest:latest
Deleted: 6dbb9cc54074106d46d4ccb330f2a40a682d49dda5f4844962b7dce9fe44aaec
```

10. Pull the image from the local registry. It should report the download of a newer image.

```
student@cp:~/app1$ sudo podman pull $repo/tagtest
```

```
Trying to pull 10.97.40.62:5000/tagtest:latest...
Getting image source signatures
Copying blob 74782b667c7d done
Copying config 6dbb9cc540 done
Writing manifest to image destination
Storing signatures
6dbb9cc54074106d46d4ccb330f2a40a682d49dda5f4844962b7dce9fe44aaec
```

11. Configure the worker (second) node to use the registry running on the cp server. Connect to the worker node. Edit the `registries.conf` file with the same values as the cp node and restart the **crio** service. Ensure it is active. You may need to install **vim** if not found.

```
student@worker:~$ sudo vim /etc/containers/registries.conf

student@worker:~$ sudo systemctl restart crio

student@worker:~$ export repo=10.97.40.62:5000

student@worker:~$ echo "export repo=10.97.40.62:5000" >> $HOME/.bashrc

student@worker:~$ sudo podman pull $repo/tagtest
```

```
Trying to pull 10.97.40.62:5000/tagtest:latest...
Getting image source signatures
Getting image source signatures
Copying blob 74782b667c7d done
Copying config 6dbb9cc540 done
Writing manifest to image destination
Storing signatures
a313e647ae05302fc54b57065c250ff58d2f580fab7b7625949651b141d9fca7
```

12. Now that we know **podman** on all nodes can use the repository we need to make sure Kubernetes knows about the new repository and settings as well. The simplest way is to reboot every node. Log back in after the connection closes.

```
student@cp:~$ sudo reboot

student@worker:~$ sudo reboot
```

13. Test that the repo works after the reboot, as good admins we want to know our configuration is persistent. Be aware it can take a minute or two after reboot for the kube-apiserver to fully start.

```
student@cp:~$ curl $repo/v2/_catalog
```

```
{"repositories": []}
```

14. Use **podman tag** to assign the `simpleapp` image and then push it to the local registry. The image and dependent images should be pushed to the local repository.

```
student@cp:~/app1$ sudo podman tag simpleapp $repo/simpleapp

student@cp:~/app1$ sudo podman push $repo/simpleapp
```

```
Getting image source signatures
Copying blob 47458fb45d99 done
Copying blob a3c1026c6bcc done
Copying blob f1d420c2af1a done
Copying blob 461719022993 done
Copying blob d35c5bda4793 done
Copying blob 46829331b1e4 done
Copying blob ceee8816bb96 done
Copying blob da7b0a80a4f2 done
Copying blob e571d2d3c73c done
Copying blob 5c2db76bc949 done
Copying config a313e647ae done
Writing manifest to image destination
Storing signatures
```

15. Test that the image can be found in the repository, from both the cp and the worker

```
student@cp:~$ curl $repo/v2/_catalog
```

```
{"repositories":["simpleapp"]}
```

16. Return to the `cp` node and deploy the `simpleapp` in Kubernetes with several replicas. We will name the deployment `try1`. Scale to have six replicas. Increase the replica count until pods are deployed on both `cp` and `worker`.

```
student@cp:~$ kubectl create deployment try1 --image=$repo/simpleapp
```

```
deployment.apps/try1 created
```

```
student@cp:~$ kubectl scale deployment try1 --replicas=6
```

```
deployment.apps/try1 scaled
```

```
student@cp:~$ kubectl get pod -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
basicpod	1/1	Running	0	20m	192.168.95.15	worker
registry-ddf6bcb7c-b8wcr	1/1	Running	0	16m	192.168.95.16	worker
try1-55f675ddd-28vgs	1/1	Running	0	17s	192.168.95.30	worker
try1-55f675ddd-2nrsj	1/1	Running	0	17s	192.168.95.26	worker
try1-55f675ddd-4vzt7	1/1	Running	0	17s	192.168.219.100	cp
<output_omitted>							

17. On the second node use **sudo crictl ps** to verify containers of `simpleapp` are running. Even though **podman** has the options, `cri-o` is running the containers on our behalf. The scheduler will usually balance pod count across nodes. As the `cp` already has several pods running the new pods may be on the `worker`, so the number of pods returned will vary.

You may need to configure **crictl** first to avoid error messages. You can omit the backslash and type the command on one line.

```
student@worker:~$ sudo crictl config \
--set runtime-endpoint=unix:///run/containerd/containerd.sock \
--set image-endpoint=unix:///run/containerd/containerd.sock
```

```
student@worker:~$ sudo crictl ps | grep simple
```

855750df82cd4	10.97.177.111:5000/					
simpleapp@sha256:f7fc297b3c20c47c506896ee13f4d7f13b4c4c0ee0fafd2ab88c9809b2a1aed6		15 minutes ago				
Running	simpleapp	0	6e31ba7491117			
237af1260265a	10.97.177.111:5000/					
simpleapp@sha256:f7fc297b3c20c47c506896ee13f4d7f13b4c4c0ee0fafd2ab88c9809b2a1aed6		15 minutes ago				
Running	simpleapp	0	bc0d63564cf49			
d0ae2910cf8b8	10.97.177.111:5000/					
simpleapp@sha256:f7fc297b3c20c47c506896ee13f4d7f13b4c4c0ee0fafd2ab88c9809b2a1aed6		15 minutes ago				
Running	simpleapp	0	b9bb90e04bec1			
046a0c439fa43	10.97.177.111:5000/					
simpleapp@sha256:f7fc297b3c20c47c506896ee13f4d7f13b4c4c0ee0fafd2ab88c9809b2a1aed6		15 minutes ago				
Running	simpleapp	0	fe464425e41b1			

18. Return to the `cp` node. Save the `try1` deployment as YAML.

```
student@cp:~/app1$ cd $HOME/app1/
student@cp:~/app1$ kubectl get deployment try1 -o yaml > simpleapp.yaml
```

19. Delete and recreate the `try1` deployment using the YAML file. Verify the deployment is running with the expected six replicas.

```
student@cp:~$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp:~/app1$ kubectl create -f simpleapp.yaml
```

```
deployment.apps/try1 created
```

```
student@cp:~/app1$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx	1/1	1	1	15m
registry	1/1	1	1	15m
try1	6/6	6	6	5s

✎ Exercise 3.3: Configure Probes

When large datasets need to be loaded or a complex application launched prior to client access, a `readinessProbe` can be used. The pod will not become available to the cluster until a test is met and returns a successful exit code. Both `readinessProbes` and `livenessProbes` use the same syntax and are identical other than the name. Where the `readinessProbe` is checked prior to being ready, then not again, the `livenessProbe` continues to be checked.

There are three types of liveness probes: a command returns a zero exit value, meaning success, an HTTP request returns a response code in the 200 to 399 range, and the third probe uses a TCP socket. In this example we'll use a command, `cat`, which will return a zero exit code when the file `/tmp/healthy` has been created and can be accessed.

1. Edit the YAML deployment file and add the stanza for a `readinessprobe`. Remember that when working with YAML whitespace matters. Indentation is used to parse where information should be associated within the stanza and the entire file. Do not use tabs. If you get an error about validating data, check the indentation. It can also be helpful to paste the file to this website to see how indentation affects the JSON value, which is actually what Kubernetes ingests: <https://www.json2yaml.com/> An edited file is also included in the tarball, but requires the image name to be edited to match your registry IP address.

```
student@cp:~/app1$ vim simpleapp.yaml
```

YAML

simpleapp.yaml

```

1  ....
2      spec:
3          containers:
4              - image: 10.111.235.60:5000/simpleapp
5                imagePullPolicy: Always
6                name: simpleapp
7                readinessProbe:           #<--This line and next five
8                  periodSeconds: 5
9                  exec:
10                     command:
11                         - cat
12                         - /tmp/healthy
13                resources: {}
14  ....
```

2. Delete and recreate the `try1` deployment.

```
student@cp:~/app1$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp:~/app1$ kubectl create -f simpleapp.yaml
```

```
deployment.apps/try1 created
```

3. The new try1 deployment should reference six pods, but show zero available. They are all missing the `/tmp/healthy` file.

```
student@cp:~/app1$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx	1/1	1	1	19m
registry	1/1	1	1	19m
try1	0/6	6	0	15s

4. Take a closer look at the pods. Use **describe pod** and **logs** to investigate issues, note there may be no logs. Choose one of the try1 pods as a test to create the health check file.

```
student@cp:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	40m
registry-795c6c8b8f-7vwdn	1/1	Running	1	40m
try1-9869bdb88-2wfnr	0/1	Running	0	26s
try1-9869bdb88-6bkn1	0/1	Running	0	26s
try1-9869bdb88-786v8	0/1	Running	0	26s
try1-9869bdb88-gmvs4	0/1	Running	0	26s
try1-9869bdb88-lfv1x	0/1	Running	0	26s
try1-9869bdb88-rtchc	0/1	Running	0	26s

5. Run the bash shell interactively and touch the `/tmp/healthy` file.

```
student@cp:~/app1$ kubectl exec -it try1-9869bdb88-rtchc -- /bin/bash
```

```
root@try1-9869bdb88-rtchc:/# touch /tmp/healthy
```

```
root@try1-9869bdb88-rtchc:/# exit
```

```
exit
```

6. Wait at least five seconds, then check the pods again. Once the probe runs again the container should show available quickly. The pod with the existing `/tmp/healthy` file should be running and show 1/1 in a READY state. The rest will continue to show 0/1.

```
student@cp:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	44m
registry-795c6c8b8f-7vwdn	1/1	Running	1	44m
try1-9869bdb88-2wfnr	0/1	Running	0	4m
try1-9869bdb88-6bkn1	0/1	Running	0	4m
try1-9869bdb88-786v8	0/1	Running	0	4m
try1-9869bdb88-gmvs4	0/1	Running	0	4m
try1-9869bdb88-lfv1x	0/1	Running	0	4m
try1-9869bdb88-rtchc	1/1	Running	0	4m

7. Touch the file in the remaining pods. Consider using a **for** loop, as an easy method to update each pod. Note the `>` shown in the output represents the secondary prompt, you would not type in that character

```
student@cp:~$ for name in try1-9869bdb88-2wfnr try1-9869bdb88-6bkn1 \
> try1-9869bdb88-786v8 try1-9869bdb88-gmvs4 try1-9869bdb88-lfvlx
> do
> kubectl exec $name -- touch /tmp/healthy
> done
```

8. It may take a short while for the probes to check for the file and the health checks to succeed.

```
student@cp:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	1h
registry-795c6c8b8f-7vwdn	1/1	Running	1	1h
try1-9869bdb88-2wfnr	1/1	Running	0	22m
try1-9869bdb88-6bkn1	1/1	Running	0	22m
try1-9869bdb88-786v8	1/1	Running	0	22m
try1-9869bdb88-gmvs4	1/1	Running	0	22m
try1-9869bdb88-lfvlx	1/1	Running	0	22m
try1-9869bdb88-rtchc	1/1	Running	0	22m

9. Now that we know when a pod is healthy, we may want to keep track that it stays healthy, using a livenessProbe. You could use one probe to determine when a pod becomes available and a second probe, to a different location, to ensure ongoing health.

Edit the deployment again. Add in a livenessProbe section as seen below. This time we will add a Sidecar container to the pod running a simple application which will respond to port 8080. Note that the dash (-) in front of the name. Also goproxy is indented the same number of spaces as the - in front of the image: line for simpleapp earlier in the file. In this example that would be seven spaces

```
student@cp:~/app1$ vim simpleapp.yaml
```

YAML

simpleapp.yaml

```
1 ....
2     terminationMessagePath: /dev/termination-log
3     terminationMessagePolicy: File
4     - name: goproxy                                #<-- Indented 7 spaces, add lines from here...
5       image: registry.k8s.io/goproxy:0.1
6       ports:
7         - containerPort: 8080
8       readinessProbe:
9         tcpSocket:
10          port: 8080
11          initialDelaySeconds: 5
12          periodSeconds: 10
13       livenessProbe:                                #<-- This line is 9 spaces indented, fyi
14         tcpSocket:
15          port: 8080
16          initialDelaySeconds: 15
17          periodSeconds: 20                          #<-- ....to here
18       dnsPolicy: ClusterFirst
19       restartPolicy: Always
20 ....
```

10. Delete and recreate the deployment.

```
student@cp:~$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp:~$ kubectl create -f simpleapp.yaml
```

```
deployment.apps/try1 created
```

11. View the newly created pods. You'll note that there are two containers per pod, and only one is running. The new `simpleapp` containers will not have the `/tmp/healthy` file, so they will not become available until we touch the `/tmp/healthy` file again. We could include a command which creates the file into the container arguments. The output below shows it can take a bit for the old pods to terminate.

```
student@cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	13h
registry-795c6c8b8f-7vwdn	1/1	Running	1	13h
try1-76cc5ffcc6-4rjvh	1/2	Running	0	3s
try1-76cc5ffcc6-bk5f5	1/2	Running	0	3s
try1-76cc5ffcc6-d8n5q	0/2	ContainerCreating	0	3s
try1-76cc5ffcc6-mm6tw	1/2	Running	0	3s
try1-76cc5ffcc6-r9q5n	1/2	Running	0	3s
try1-76cc5ffcc6-tx4dz	1/2	Running	0	3s
try1-9869bdb88-2wfnr	1/1	Terminating	0	12h
try1-9869bdb88-6bknl	1/1	Terminating	0	12h
try1-9869bdb88-786v8	1/1	Terminating	0	12h
try1-9869bdb88-gmvs4	1/1	Terminating	0	12h
try1-9869bdb88-lfv1x	1/1	Terminating	0	12h
try1-9869bdb88-rtchc	1/1	Terminating	0	12h

12. Create the health check file for the readinessProbe. You can use a **for** loop again for each action, this setup will leverage labels so you don't have to look up the pod names. As there are now two containers in the pod, you should include the container name for which one will execute the command. If no name is given, it will default to the first container. Depending on how you edited the YAML file `try1` should be the first pod and `goproxy` the second. To ensure the correct container is updated, add **-c simpleapp** to the `kubectl` command. Your pod names will be different. Use the names of the newly started containers from the `kubectl get pods` command output. Note the `>` character represents the secondary prompt, you would not type in that character.

```
student@cp:~$ for name in $(kubectl get pod -l app=try1 -o name)
> do
> kubectl exec $name -c simpleapp -- touch /tmp/healthy
> done
```

13. In the next minute or so the Sidecar container in each pod, which was not running, will change status to Running. Each should show 2/2 containers running.

```
student@cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-g7lnk	1/1	Running	1	13h
registry-795c6c8b8f-7vwdn	1/1	Running	1	13h
try1-76cc5ffcc6-4rjvh	2/2	Running	0	3s
try1-76cc5ffcc6-bk5f5	2/2	Running	0	3s
try1-76cc5ffcc6-d8n5q	2/2	Running	0	3s
try1-76cc5ffcc6-mm6tw	2/2	Running	0	3s
try1-76cc5ffcc6-r9q5n	2/2	Running	0	3s
try1-76cc5ffcc6-tx4dz	2/2	Running	0	3s

14. View the events for a particular pod. Even though both containers are currently running and the pod is in good shape, note the events section shows the issue, but not a change in status or the probe success.

```
student@cp:~/app1$ kubectl describe pod try1-76cc5ffcc6-tx4dz | tail
```



```

Events:
  Type    Reason      Age           From              Message
  ----    -
Normal    Scheduled   7m46s         default-scheduler  Successfully assigned
                        default/try1-754bf9c75b-vx58x to cp
Normal    Pulling     7m44s         kubelet            Pulling image
                        "10.97.177.111:5000/simpleapp"
Normal    Pulled      7m44s         kubelet            Successfully pulled image
                        "10.97.177.111:5000/simpleapp" in 96.710062ms
Normal    Created     7m43s         kubelet            Created container simpleapp
Normal    Started     7m43s         kubelet            Started container simpleapp
Warning   Unhealthy   7m6s (x9 over 7m42s) kubelet            Readiness probe failed:
                        cat: /tmp/healthy: No such file or directory

```

15. If you look for the status of each container in the pod, they should show that both are Running and ready showing True.

```
student@cp:~/app1$ kubectl describe pod try1-76cc5ffcc6-tx4dz | grep -E 'State|Ready'
```

```

State:      Running
Ready:      True
State:      Running
Ready:      True
Ready:      True
ContainersReady  True

```

Exercise 3.4: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

- Implement probes and health checks
- Define, build and modify container images
- Understand multi-container Pod design patterns
- Utilize container logs

Focus on ensuing you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of YAML samples and can complete each step quickly.

The domain review items may be vague or non-detailed, just like what one would experience in a work or exam environment.

1. Using the three URL locations allowed by the exam, find and bookmark working YAML examples for LivenessProbes, ReadinessProbes, and multi-container pods.
2. Deploy a new nginx webserver. Add a LivenessProbe and a ReadinessProbe on port 80. Test that both probes and the webserver work.
3. Use the `build-review1.yaml` file to create a non-working deployment. Fix the deployment such that both containers are running and in a READY state. The web server listens on port 80, and the proxy listens on port 8080.

4. View the default page of the web server. When successful verify the GET activity logs in the container log. The message should look something like the following. Your time and IP may be different.

```
192.168.124.0 - - [30/Jan/2020:03:30:31 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.58.0" "-"
```

5. Remove any resources created in this review.

Chapter 4

Design



Exercise 4.1: Planning the Deployment

Overview

In this exercise we will investigate common network plugins. Each **kubelet** agent uses one plugin at a time. Due to complexity, the entire cluster uses one plugin which is configured prior to application deployment. Some plugins don't honor security configurations such as network policies. Should you design a deployment which and use a network policy there wouldn't be an error; the policy would have no effect. While developers typically wouldn't care much about the mechanics of it can affect the availability of features and troubleshooting of newly decoupled microservices.

While still new, the community is moving towards the **Container Network Interface (CNI)** specification (<https://github.com/containernetworking/cni>). This provides the most flexibility and features in the fast changing space of container networking.

A common alternative is **kubenet**, a basic plugin which relies on the cloud provider to handle routing and cross-node networking. In a previous lab exercise we configured **Project Calico**. Classic and external modes are also possible. Several software defined network projects intended for Kubernetes have been created recently, with new features added regularly. Be aware that **Calico** is a dynamic project with ongoing and frequent changes.

Quick Understanding of Network Plugins

While developers don't need to configure cluster networking, they may need to understand the plugin in use, and its particular features and quirks. This section is to ensure you have made a quick review of the most commonly used plugins, and where to find more information as necessary.

1. Verify your nodes are using a CNI plugin. Read through the startup process of CNI. Each message begins with a type of message and a time stamp, followed by the details of the message. Use **TAB** to complete for your node name. Examine both the controller and one of your proxy pods.

```
student@cp:~$ kubectl -n kube-system logs kube-controller-manager-<TAB>
```

```
I0216 17:06:11.729548      1 serving.go:348] Generated self-signed cert in-memory
I0216 17:06:12.318263      1 controllermanager.go:196] Version: v1.25.1
I0216 17:06:12.321279      1 dynamic_cafire_content.go:156] "Starting controller" name=....
I0216 17:06:12.321311      1 secure_serving.go:200] Serving securely on 127.0.0.1:10257
....
```

```
student@cp:~$ kubectl -n kube-system logs kube-proxy-<TAB>
```

```
I0216 17:06:19.299029      1 node.go:163] Successfully retrieved node IP: 10.128.0.67
I0216 17:06:19.299337      1 server_others.go:138] "Detected node IP" address="10.128.0.67"
I0216 17:06:19.299515      1 server_others.go:561] "Unknown proxy mode, assuming iptables proxy"
↪ proxyMode=""
```

2. There are many CNI providers possible. The following list represents some of the more common choices, but it is not exhaustive. With many new plugins being developed there may be another which better serves your needs. Use these websites to answer questions which follow. While we strive to keep the answers accurate, please be aware that this area has a lot of attention and development and changes often.

- **Project Calico**

<https://docs.projectcalico.org/v3.0/introduction/>

- **Calico with Canal**

<https://docs.projectcalico.org/v3.0/getting-started/kubernetes/installation/hosted/canal>

- **Weave Works**

<https://www.weave.works/docs/net/latest/kubernetes/kube-addon>

- **Flannel**

<https://github.com/coreos/flannel>

- **Cilium**

<http://cilium.io/>

- **Kube Router**

<https://www.kube-router.io>

3. Which of the plugins allow vxlans?
4. Which are layer 2 plugins?
5. Which are layer 3?
6. Which allow network policies?
7. Which can encrypt all TCP and UDP traffic?

Multi-container Pod Considerations

Using the information learned from this chapter, consider the following questions:

1. Which deployment method would allow the most flexibility, multiple applications per pod or one per pod?
2. Which deployment method allows for the most granular scalability?
3. Which have the best performance?
4. How many IP addresses are assigned per pod?
5. What are some ways containers can communicate within the same pod?
6. What are some reasons you should have multiple containers per pod?

Do you really know?

When and why would you use a multi-container pod?

Have you found a YAML example online?

Go back and review multi-container pod types and content on decoupling if you can't easily answer these questions. We touched on adding a second logging and a readiness container in a previous chapter and will work more with logging a future exercise.

✔ Solution 4.1

Plugin Answers

1. Which of the plugins allow vxlans?
Canal, Project Calico, Flannel, Weave Net, Cilium
2. Which are layer 2 plugins?
Canal, Flannel, Weave Net
3. Which are layer 3?
Project Calico, Romana, Kube Router
4. Which allow network policies?
Project Calico, Canal, Kube Router, Weave Net, Cilium
5. Which can encrypt all TCP and UDP traffic?
Project Calico, Weave Net, Cilium

Multi Pod Answers

1. Which deployment method would allow the most flexibility, multiple applications per pod or one per Pod?
One per pod
2. Which deployment method allows for the most granular scalability?
One per pod
3. Which have the best inter-container performance?
Multiple per pod.
4. How many IP addresses are assigned per pod?
One
5. What are some ways containers can communicate within the same pod?
IPC, loopback or shared filesystem access.
6. What are some reasons you should have multiple containers per pod?
Lean containers may not have functionality like logging. Able to maintain lean execution but add functionality as necessary, like Ambassadors and Sidecar containers.

✍ Exercise 4.2: Designing Applications With Duration: Create a Job

While most applications are deployed such that they continue to be available there are some which we may want to run a particular number of times called a Job, and others on a regular basis called a CronJob

1. Create a job which will run a container which sleeps for three seconds then stops.

```
student@cp:~$ vim job.yaml
```

YAML

job.yaml

```
1 apiVersion: batch/v1
2 kind: Job
3 metadata:
4   name: sleepy
5 spec:
6   template:
7     spec:
8       containers:
9         - name: resting
10           image: busybox
11           command: ["/bin/sleep"]
12           args: ["3"]
13           restartPolicy: Never
```

2. Create the job, then verify and view the details. The example shows checking the job three seconds in and then again after it has completed. You may see different output depending on how fast you type.

```
student@cp:~$ kubectl create -f job.yaml
```

```
job.batch/sleepy created
```

```
student@cp:~$ kubectl get job
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	0/1	3s	3s

```
student@cp:~$ kubectl describe jobs.batch sleepy
```

```
Name:          sleepy
Namespace:     default
Selector:      controller-uid=24c91245-d0fb-11e8-947a-42010a800002
Labels:        controller-uid=24c91245-d0fb-11e8-947a-42010a800002
               job-name=sleepy
Annotations:   <none>
Parallelism:   1
Completions:   1
Start Time:    Sun, 03 Nov 2019 04:22:50 +0000
Completed At:  Sun, 03 Nov 2019 04:22:55 +0000
Duration:      5s
Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
<output_omitted>
```

```
student@cp:~$ kubectl get job
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	1/1	5s	17s

3. View the configuration information of the job. There are three parameters we can use to affect how the job runs. Use `-o yaml` to see these parameters. We can see that `backoffLimit`, `completions`, and the `parallelism`. We'll add these parameters next.

```
student@cp:~$ kubectl get jobs.batch sleepy -o yaml
```

```
<output_omitted>
  uid: c2c3a80d-d0fc-11e8-947a-42010a800002
  spec:
    backoffLimit: 6
    completions: 1
    parallelism: 1
    selector:
      matchLabels:
<output_omitted>
```

4. As the job continues to AGE in a completion state, delete the job.

```
student@cp:~$ kubectl delete jobs.batch sleepy
```

```
job.batch "sleepy" deleted
```

5. Edit the YAML and add the `completions:` parameter and set it to 5.

```
student@cp:~$ vim job.yaml
```

YAML

job.yaml

```
1 <output_omitted>
2 metadata:
3   name: sleepy
4 spec:
5   completions: 5  #<--Add this line
6   template:
7     spec:
8       containers:
9 <output_omitted>
```

6. Create the job again. As you view the job note that COMPLETIONS begins as zero of 5.

```
student@cp:~$ kubectl create -f job.yaml
```

```
job.batch/sleepy created
```

```
student@cp:~$ kubectl get jobs.batch
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	0/5	5s	5s

7. View the pods that running. Again the output may be different depending on the speed of typing.

```
student@cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-67f8fb575f-g4468	1/1	Running	2	2d
registry-56cffc98d6-xlhhf	1/1	Running	1	2d
sleepy-z5tnh	0/1	Completed	0	8s
sleepy-zd692	1/1	Running	0	3s

<output_omitted>

8. Eventually all the jobs will have completed. Verify then delete the job.

```
student@cp:~$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	5/5	26s	10m

```
student@cp:~$ kubectl delete jobs.batch sleepy
```

```
job.batch "sleepy" deleted
```

9. Edit the YAML again. This time add in the `parallelism:` parameter. Set it to 2 such that two pods at a time will be deployed.

```
student@cp:~$ vim job.yaml
```

YAML

job.yaml

```
1 <output_omitted>
2   name: sleepy
3   spec:
4     completions: 5
5     parallelism: 2    #<-- Add this line
6     template:
7       spec:
8 <output_omitted>
```

10. Create the job again. You should see the pods deployed two at a time until all five have completed.

```
student@cp:~$ kubectl create -f job.yaml
```

```
student@cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-67f8fb575f-g4468	1/1	Running	2	2d
registry-56cffc98d6-xlhbf	1/1	Running	1	2d
sleepy-8xwpc	1/1	Running	0	5s
sleepy-xjqnf	1/1	Running	0	5s
try1-c9cb54f5d-b45gl	2/2	Running	0	8h
<output_omitted>				

```
student@cp:~$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	3/5	11s	11s

11. Add a parameter which will stop the job after a certain number of seconds. Set the `activeDeadlineSeconds:` to 15. The job and all pods will end once it runs for 15 seconds.

```
student@cp:~$ vim job.yaml
```

YAML

job.yaml

```
1 <output_omitted>
2   completions: 5
3   parallelism: 2
4   activeDeadlineSeconds: 15    #<-- Add this line
5   template:
6     spec:
7       containers:
```




```

8     - name: resting
9       image: busybox
10      command: ["/bin/sleep"]
11      args: ["3"]
12 <output_omitted>

```

12. Delete and recreate the job again. It should run for four times then continue to age without further completions.

```
student@cp:~$ kubectl delete jobs.batch sleepy
```

```
job.batch "sleepy" deleted
```

```
student@cp:~$ kubectl create -f job.yaml
```

```
job.batch/sleepy created
```

```
student@cp:~$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	2/5	6s	6s

```
student@cp:~$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy	4/5	16s	16s

13. View the message: entry in the Status section of the object YAML output. You may see less status if the job has yet to run. Wait and try again, if so.

```
student@cp:~$ kubectl get job sleepy -o yaml
```

```

<output_omitted>
status:
  conditions:
  - lastProbeTime: "2019-11-03T16:06:10Z"
    lastTransitionTime: "2019-11-03T16:06:10Z"
    message: Job was active longer than specified deadline
    reason: DeadlineExceeded
    status: "True"
    type: Failed
  failed: 1
  startTime: "2019-11-03T16:05:55Z"
  succeeded: 4

```

14. Delete the job.

```
student@cp:~$ kubectl delete jobs.batch sleepy
```

```
job.batch "sleepy" deleted
```

Exercise 4.3: Designing Applications With Duration: Create a CronJob

A CronJob creates a watch loop which will create a batch job on your behalf when the time becomes true. We will use our existing Job file to start.

1. Copy the Job file to a new file.

```
student@cp:~$ cp job.yaml cronjob.yaml
```

2. Edit the file to look like the annotated file shown below.

```
student@cp:~$ vim cronjob.yaml
```

YAML

cronjob.yaml

```
1 apiVersion: batch/v1
2 kind: CronJob    #<-- Change this line
3 metadata:
4   name: sleepy
5 spec:
6   schedule: "*/2 * * * *"    #<-- Add Linux style cronjob syntax
7   jobTemplate:               #<-- New jobTemplate and spec
8     spec:
9       template: #<-- This and following lines space four to right
10         spec:
11           containers:
12             - name: resting
13               image: busybox
14               command: ["/bin/sleep"]
15               args: ["3"]
16             restartPolicy: Never
```

3. Create the new CronJob. View the jobs. It will take two minutes for the CronJob to run and generate a new batch Job.

```
student@cp:~$ kubectl create -f cronjob.yaml
```

```
cronjob.batch/sleepy created
```

```
student@cp:~$ kubectl get cronjobs.batch
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
sleepy	*/2 * * * *	False	0	<none>	8s

```
student@cp:~$ kubectl get job
```

```
No resources found in default namespace.
```

4. After two minutes you should see jobs start to run.

```
student@cp:~$ kubectl get cronjobs.batch
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
sleepy	*/2 * * * *	False	0	21s	2m1s

```
student@cp:~$ kubectl get jobs.batch
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539722040	1/1	5s	18s

```
student@cp:~$ kubectl get jobs.batch
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539722040	1/1	5s	5m17s
sleepy-1539722160	1/1	6s	3m17s
sleepy-1539722280	1/1	6s	77s

5. Ensure that if the job continues for more than 10 seconds it is terminated. We will first edit the **sleep** command to run for 30 seconds then add the `activeDeadlineSeconds`: entry to the container.

```
student@cp:~$ vim cronjob.yaml
```

YAML

cronjob.yaml

```
1 ....
2   jobTemplate:
3     spec:
4       template:
5         spec:
6           activeDeadlineSeconds: 10 #<-- Add this line
7           containers:
8             - name: resting
9             ....
10            command: ["/bin/sleep"]
11            args: ["30"] #<-- Edit this line
12            restartPolicy: Never
```

6. Delete and recreate the CronJob. It may take a couple of minutes for the batch Job to be created and terminate due to the timer.

```
student@cp:~$ kubectl delete cronjobs.batch sleepy
```

```
cronjob.batch "sleepy" deleted
```

```
student@cp:~$ kubectl create -f cronjob.yaml
```

```
cronjob.batch/sleepy created
```

```
student@cp:~$ sleep 120 ; kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539723240	0/1	61s	61s

```
student@cp:~$ kubectl get cronjobs.batch
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
sleepy	*/* * * * *	False	1	72s	94s

```
student@cp:~$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539723240	0/1	75s	75s

```
student@cp:~$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
sleepy-1539723240	0/1	2m19s	2m19s
sleepy-1539723360	0/1	19s	19s

```
student@cp:~$ kubectl get cronjobs.batch
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
sleepy	*/* * * * *	False	2	31s	2m53s

7. Clean up by deleting the CronJob.

```
student@cp:~$ kubectl delete cronjobs.batch sleepy
```

```
cronjob.batch "sleepy" deleted
```

Exercise 4.4: Using Labels

Create and work with labels. We will understand how the deployment, replicaSet, and pod labels interact.

1. Create a new deployment called design2

```
student@cp:~$ kubectl create deployment design2 --image=nginx
```

```
deployment.apps/design2 created
```

2. View the wide **kubectl get** output for the design2 deployment and make note of the SELECTOR

```
student@cp:~$ kubectl get deployments.apps design2 -o wide
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
design2	1/1	1	1	2m13s	nginx	nginx	app=design2

3. Use the **-l** option to use the selector to list the pods running inside the deployment. There should be only one pod running.

```
student@cp:~$ kubectl get -l app=design2 pod
```

NAME	READY	STATUS	RESTARTS	AGE
design2-766d48574f-5w274	1/1	Running	0	3m1s

4. View the pod details in YAML format using the deployment selector. This time use the **--selector** option. Find the pod label in the output. It should match that of the deployment.

```
student@cp:~$ kubectl get --selector app=design2 pod -o yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    cni.projectcalico.org/podIP: 192.168.113.222/32
    creationTimestamp: "2020-01-31T16:29:37Z"
    generateName: design2-766d48574f-
  labels:
    app: design2
    pod-template-hash: 766d48574f
....
```

5. Edit the pod label to be your favorite color.

```
student@cp:~$ kubectl edit pod design2-766d48574f-5w274
```

YAML

```

1 .....
2   labels:
3       app: orange                                #<<-- Edit this line
4       pod-template-hash: 766d48574f
5       name: design2-766d48574f-5w274
6   ....

```

6. Now view how many pods are in the deployment. Then how many have design2 in their name. Note the AGE of the pods.

```
student@cp:~$ kubectl get deployments.apps design2 -o wide
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
design2	1/1	1	1	56s	nginx	nginx	app=design2

```
student@cp:~$ kubectl get pods | grep design2
```

design2-766d48574f-5w274	1/1	Running	0	82s
design2-766d48574f-xttgg	1/1	Running	0	2m12s

7. Delete the design2 deployment.

```
student@cp:~$ kubectl delete deploy design2
```

```
deployment.apps "design2" deleted
```

8. Check again for pods with design2 in their names. You should find one pod, with an AGE of when you first created the deployment. Once the label was edited the deployment created a new pod in order that the status matches the spec and there be a replica running with the intended label.

```
student@cp:~$ kubectl get pods | grep design2
```

design2-766d48574f-5w274	1/1	Running	0	38m
--------------------------	-----	---------	---	-----

9. Delete the pod using the **-l** and the label you edited to be your favorite color in a previous step. The command details have been omitted. Use previous steps to figure out these commands.

Exercise 4.5: Setting Pod Resource Limits and Requirements

1. Create a new pod running the vish/stress image. A YAML `stress.yaml` file has been included in the course tarball.
2. Run the **top** command on the cp and worker nodes. You should find a stress command consuming the majority of the CPU on one node, the worker. Use **q** to exit from top. Delete the deployment.
3. Edit the `stress.yaml` file add in the following limits and requests.

```
student@cp:~$
```

YAML

```

1 .....
2   name: stressmeout
3   resources:                                #<<-- Add this and following six lines
4       limits:
5           cpu: "1"
6           memory: "1Gi"
7       requests:
8           cpu: "0.5"

```

YAML

```

9       memory: "500Mi"
10     args:
11       - -cpus
12     ....

```

4. Create the deployment again. Check the status of the pod. You should see that it shows an OOMKilled status and a growing number of restarts. You may see a status of Running if you catch the pod in early in a restart. If you wait long enough you may see CrashLoopBackOff.

```
student@cp:~$ kubectl get pod stressmeout-7fbbbcc887-v9kvb
```

NAME	READY	STATUS	RESTARTS	AGE
stressmeout-7fbbbcc887-v9kvb	0/1	OOMKilled	2	32s

5. Delete then edit the deployment. Change the `limit:` parameters such that pod is able to run, but not too much extra resources. Try setting the memory limit to exactly what the stress command requests.

As we allow the pod to run on the cp node, this could cause issues, such as the kube-apiserver restarting due to lack of resources. We will also add a `nodeSelector` to use the built in label of `kubernetes.io/hostname`.

```
student@cp:~$ kubectl delete -f stress.yaml
```

```
student@cp:~$ vim stress.yaml
```

YAML

```

1     ....
2   spec:
3     nodeSelector:                #<-- Uncomment and edit
4       kubernetes.io/hostname: worker #<-- to by YOUR worker hostname
5     containers:
6
7     ....
8     resources:
9       limits:
10        cpu: "2"
11        memory: "2Gi"
12      requests:
13      ....

```

6. Create the deployment and ensure the pod runs without error. Use **top** to verify the stress command is running on one of the nodes and view the pod details to ensure the CPU and memory limits are in use. Also use the **kubectl describe node** command to view the resources your cp and worker node are using. The command details have been omitted. Use previous steps to figure out the commands.
7. Change `limits` and `requests` to numbers higher than your node resources, and evaluate how the container and pod is handled. It may take a while for resources to be fully allocated.
8. Remove the stressmeout deployment when done.

```
student@cp:~$ kubectl delete deploy stressmeout
```

```
deployment.apps "stressmeout" deleted
```

✍ Exercise 4.6: Simple initContainer

1. As we have already learned about creating pods, this exercise does not include the particular steps. You can use this to determine if you need to review previous content. Use the following YAML, which is also in the course tarball, to create

a pod with an `initContainer`. It will fail, view what it looks like and investigate the errors. Change the command from `/bin/false` to a successful command such as `/bin/true` and create the pod again.



init-tester.yaml

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: init-tester
5   labels:
6     app: inittest
7 spec:
8   containers:
9     - name: webservice
10      image: nginx
11   initContainers:
12     - name: failed
13      image: busybox
14      command: [/bin/false]
```

✍ Exercise 4.7: Exploring Custom Resource Definitions

1. View CRDs currently in the cluster

```
student@cp:~$ kubectl get crd
```

AME	CREATED AT
bgpconfigurations.crd.projectcalico.org	2021-10-13T04:57:12Z
bgppeers.crd.projectcalico.org	2021-10-13T04:57:12Z
blockaffinities.crd.projectcalico.org	2021-10-13T04:57:12Z
clusterinformations.crd.projectcalico.org	2021-10-13T04:57:12Z
felixconfigurations.crd.projectcalico.org	2021-10-13T04:57:12Z
globalnetworkpolicies.crd.projectcalico.org	2021-10-13T04:57:13Z
globalnetworksets.crd.projectcalico.org	2021-10-13T04:57:13Z
hostendpoints.crd.projectcalico.org	2021-10-13T04:57:13Z
ipamblocks.crd.projectcalico.org	2021-10-13T04:57:13Z
ipamconfigs.crd.projectcalico.org	2021-10-13T04:57:13Z
ipamhandles.crd.projectcalico.org	2021-10-13T04:57:13Z
ippools.crd.projectcalico.org	2021-10-13T04:57:13Z
kubecontrollersconfigurations.crd.projectcalico.org	2021-10-13T04:57:14Z
networkpolicies.crd.projectcalico.org	2021-10-13T04:57:14Z
networksets.crd.projectcalico.org	2021-10-13T04:57:14Z
serviceprofiles.linkerd.io	2021-10-13T05:00:28Z
trafficsplits.split.smi-spec.io	2021-10-13T05:00:28Z

2. Take a detailed look at the IPPools CRD.

```
student@cp:~$ kubectl get crd ippools.crd.projectcalico.org -o yaml
```

```
....
spec:
  conversion:
    strategy: None
  group: crd.projectcalico.org
  names:
    kind: IPPool
    listKind: IPPoolList
    plural: ippools
    singular: ippool
  scope: Cluster
  versions:
```

```
- name: v1
....
```

3. View if this newly declared resource exists.

```
student@cp:~$ kubectl get IPPool
```

```
NAME                AGE
default-ipv4-ippool 11h
```

4. Take a detailed view of the resource. Note the IP information should match the ephemeral pod IPs your cluster is using.

```
student@cp:~$ kubectl get IPPool -o yaml
```

```
....
uid: 0640c9e5-25f0-40f8-9c36-1d9be26a2516
spec:
  blockSize: 26
  cidr: 192.168.0.0/16
  ipipMode: Always
  natOutgoing: true
  nodeSelector: all()
  vxlanMode: Never
kind: List
metadata:
....
```

Exercise 4.8: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter. They may be in multiple sections. The items below shows the topics covered in this chapter.

- Understand Jobs and CronJobs
- Understanding and defining resources requirements, limits and quotas
- Understand multi-container Pod design patterns (e.g. sidecar **init**, and others)
- Discover and use resources that extend Kubernetes (CRD).

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

1. Find and use the `design-review1.yaml` file to create a pod.
2. Determine the CPU and memory resource requirements of `design-pod1`.
3. Edit the pod resource requirements such that the CPU limit is exactly twice the amount requested by the container. (Hint: subtract .22)
4. Increase the memory resource limit of the pod until the pod shows a `Running` status. This may require multiple edits and attempts. Determine the minimum amount necessary for the `Running` status to persist at least a minute.

5. Use the `design-review2.yaml` file to create several pods with various labels.
6. Using **only** the `--selector` value `tux` to delete only those pods. This should be half of the pods. Hint, you will need to view pod settings to determine the key value as well.
7. Create a new cronjob which runs `busybox` and the `sleep 30` command. Have the cronjob run every three minutes. View the job status to check your work. Change the settings so the pod runs 10 minutes from the current time, every week. For example, if the current time was 2:14PM, I would configure the job to run at 2:24PM, every Monday.
8. Delete any objects created during this review. You may want to delete all but the cronjob if you'd like to see if it runs in 10 minutes. Then delete that object as well.

Chapter 5

Deployment Configuration



Exercise 5.1: Configure the Deployment: ConfigMaps



Very Important

Save a copy of your `$HOME/app1/simpleapp.yaml` file, in case you would like to repeat portions of the labs, or you find your file difficult to use due to typos and whitespace issues.

```
student@cp:~$ cp $HOME/app1/simpleapp.yaml $HOME/beforeLab5.yaml
```

We will cover the use of **secrets** in the Security chapter lab.

Overview

In this lab we will add resources to our deployment with further configuration you may need for production.

There are three different ways a **ConfigMap** can ingest data, from a literal value, from a file, or from a directory of files.

1. Create a **ConfigMap** containing primary colors. We will create a series of files to ingest into the **ConfigMap**. First create a directory **primary** and populate it with four files. Then we create a file in our home directory with our favorite color.

```
student@cp:~/app1$ cd

student@cp:~$ mkdir primary
student@cp:~$ echo c > primary/cyan
student@cp:~$ echo m > primary/magenta
student@cp:~$ echo y > primary/yellow
student@cp:~$ echo k > primary/black
student@cp:~$ echo "known as key" >> primary/black
student@cp:~$ echo blue > favorite
```

2. Generate a **configMap** using each of the three methods.

```
student@cp:~$ kubectl create configmap colors \
  --from-literal=text=black \
  --from-file=./favorite \
  --from-file=./primary/
```

```
configmap/colors created
```

3. View the newly created **configMap**. Note the way the ingested data is presented.

```
student@cp:~$ kubectl get configmap colors
```

NAME	DATA	AGE
colors	6	11s

```
student@cp:~$ kubectl get configmap colors -o yaml
```

```
apiVersion: v1
data:
  black: |
    k
    known as key
  cyan: |
    c
  favorite: |
    blue
  magenta: |
    m
  text: black
  yellow: |
    y
kind: ConfigMap
metadata:
  <output_omitted>
```

4. Update the YAML file of the application to make use of the **configMap** as an environmental parameter. Add the six lines from the env: line to key:favorite.

```
student@cp:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```
1 ....
2 spec:
3   containers:
4     - image: 10.105.119.236:5000/simpleapp
5       env: #<-- Add from here
6     - name: ilike
7       valueFrom:
8         configMapKeyRef:
9           name: colors
10          key: favorite #<-- to here
11     imagePullPolicy: Always
12 ....
```

5. Delete and re-create the deployment with the new parameters.

```
student@cp-lab-7xtx:~$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp-lab-7xtx:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
deployment.apps/try1 created
```

6. Even though the `try1` pod is not in a fully ready state, it is running and useful. Use **kubectl exec** to view a variable's value. View the pod state then verify you can see the `ilike` value within the `simpleapp` container. Note that the use of double dash (`--`) tells the shell to pass the following as standard in.

```
student@cp:~$ kubectl get pod
```

```
<output_omitted>
```

```
student@cp:~$ kubectl exec -c simpleapp -it try1-5db9bc6f85-whxbf \
-- /bin/bash -c 'echo $ilike'
```

```
blue
```

7. Edit the YAML file again, this time adding the another method of using a **configMap**. Edit the file to add three lines. `envFrom` should be indented the same amount as `env` earlier in the file, and `configMapRef` should be indented the same as `configMapKeyRef`.

```
student@cp:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```
1 ....
2         configMapKeyRef:
3             name: colors
4             key: favorite
5     envFrom: #<-- Add this and the following two lines
6     - configMapRef:
7         name: colors
8     imagePullPolicy: Always
9 ....
```

8. Again delete and recreate the deployment. Check the pods restart.

```
student@cp:~$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
deployment.apps/try1 created
```

```
student@cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	23h
registry-795c6c8b8f-hl5w	1/1	Running	2	23h
try1-d4fbf76fd-46pkb	1/2	Running	0	40s
try1-d4fbf76fd-9kw24	1/2	Running	0	39s
try1-d4fbf76fd-bx9j9	1/2	Running	0	39s
try1-d4fbf76fd-jw8g7	1/2	Running	0	40s
try1-d4fbf76fd-lpp15	1/2	Running	0	39s
try1-d4fbf76fd-xtfd4	1/2	Running	0	40s

9. View the settings inside the `try1` container of a pod. The following output is truncated in a few places. Omit the container name to observe the behavior. Also execute a command to see all environmental variables instead of logging into the container first.

```
student@cp:~$ kubectl exec -it try1-d4fbf76fd-46pkb -- /bin/bash -c 'env'
```

```

Defaulting container name to simpleapp.
Use 'kubectl describe pod/try1-d4fbf76fd-46pkb -n default' to see all of the containers in this
↳ pod.
REGISTRY_PORT_5000_TCP_ADDR=10.105.119.236
HOSTNAME=try1-d4fbf76fd-46pkb
TERM=xterm
yellow=y
<output_omitted>
REGISTRY_SERVICE_HOST=10.105.119.236
KUBERNETES_SERVICE_PORT=443
REGISTRY_PORT_5000_TCP=tcp://10.105.119.236:5000
KUBERNETES_SERVICE_HOST=10.96.0.1
text=black
REGISTRY_SERVICE_PORT_5000=5000
<output_omitted>
black=k
known as key

<output_omitted>
ilike=blue
<output_omitted>
magenta=m

cyan=c
<output_omitted>

```

10. For greater flexibility and scalability **ConfigMaps** can be created from a YAML file, then deployed and redeployed as necessary. Once ingested into the cluster the data can be retrieved in the same manner as any other object. Create another **configMap**, this time from a YAML file.

```
student@cp:~$ vim car-map.yaml
```

YAML

car-map.yaml

```

1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: fast-car
5   namespace: default
6 data:
7   car.make: Ford
8   car.model: Mustang
9   car.trim: Shelby

```

```
student@cp:~$ kubectl create -f car-map.yaml
```

```
configmap/fast-car created
```

11. View the ingested data, note that the output is just as in file created.

```
student@cp:~$ kubectl get configmap fast-car -o yaml
```

```

apiVersion: v1
data:
  car.make: Ford
  car.model: Mustang
  car.trim: Shelby
kind: ConfigMap
metadata:
  <output_omitted>

```

12. Add the **configMap** settings to the `simpleapp.yaml` file as a volume. Both containers in the `try1` deployment can access to the same volume, using `volumeMounts` statements. Remember that the volume stanza is of equal depth to the containers stanza, and should come after the containers have been declared, the example below has the volume added just before the `status:` output..

```
student@cp:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```
1  ....
2      spec:
3          containers:
4              - image: 10.105.119.236:5000/simpleapp
5                volumeMounts:          #<-- Add this and following two lines
6                  - mountPath: /etc/cars
7                    name: car-vol
8              env:
9                  - name: ilike
10     ....
11     securityContext: {}
12     terminationGracePeriodSeconds: 30
13     volumes:          #<-- Add this and following four lines
14         - name: car-vol
15           configMap:
16             defaultMode: 420
17             name: fast-car
18 status:
19     ....
```

13. Delete and recreate the deployment.

```
student@cp:~$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
deployment.apps/try1 created
```

14. Verify the deployment is running. Note that we still have not automated the creation of the `/tmp/healthy` file inside the container, as a result the `AVAILABLE` count remains zero until we use the **for** loop to create the file. We will remedy this in the next step.

```
student@cp:~$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx	1/1	1	1	1d
registry	1/1	1	1	1d
try1	0/6	6	0	39s

15. Our health check was the successful execution of a command. We will edit the command of the existing `readinessProbe` to check for the existence of the mounted `configMap` file and re-create the deployment. After a minute both containers should become available for each pod in the deployment. Be sure you edit the `simpleapp` section, not the `goproxy` section.

```
student@cp:~$ kubectl delete deployment try1
```

```
deployment.apps "try1" deleted
```

```
student@cp:~$ vim $HOME/app1/simpleapp.yaml
```

YAML

simpleapp.yaml

```
1  ....
2      readinessProbe:
3          exec:
4              command:
5                  - ls /etc/cars #<-- Add/Edit this and following line.
6                  - /etc/cars
7          periodSeconds: 5
8  ....
```

```
student@cp:~$ kubectl create -f $HOME/app1/simpleapp.yaml
```

```
deployment.apps/try1 created
```

16. Wait about a minute and view the deployment and pods. All six replicas should be running and report that 2/2 containers are in a ready state within.

```
student@cp:~$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx	1/1	1	1	1d
registry	1/1	1	1	1d
try1	6/6	6	6	1m

```
student@cp:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	1d
registry-795c6c8b8f-hl5wf	1/1	Running	2	1d
try1-7865dcb948-2dzc8	2/2	Running	0	1m
try1-7865dcb948-7fkh7	2/2	Running	0	1m
try1-7865dcb948-d85bc	2/2	Running	0	1m
try1-7865dcb948-djrcj	2/2	Running	0	1m
try1-7865dcb948-kwlv8	2/2	Running	0	1m
try1-7865dcb948-stb2n	2/2	Running	0	1m

17. View a file within the new volume mounted in a container. It should match the data we created inside the configMap. Because the file did not have a carriage-return it will appear prior to the following prompt.

```
student@cp:~$ kubectl exec -c simpleapp -it try1-7865dcb948-stb2n \
-- /bin/bash -c 'cat /etc/cars/car.trim'
```

```
Shelby student@cp:~$
```

✍ Exercise 5.2: Configure the Deployment: Attaching Storage

There are several types of storage which can be accessed with Kubernetes, with flexibility of storage being essential to scalability. In this exercise we will configure an NFS server. With the NFS server we will create a new **persistent volume (pv)** and a **persistent volume claim (pvc)** to use it.

1. Search for pv and pvc YAML example files on <http://kubernetes.io/docs> and <http://kubernetes.io/blog>.
2. Use the `CreateNFS.sh` script from the tarball to set up NFS on your cp node. This script will configure the server, export `/opt/sfw` and create a file `/opt/sfw/hello.txt`. Use the `find` command to locate the file if you don't remember where you extracted the tar file. This example narrows the search to your `$HOME` directory. Change for your environment. You may find the same file in more than one sub-directory of the tarfile.


```
student@cp:~$ find $HOME -name CreateNFS.sh
```

```
<some_path>/CreateNFS.sh
```

```
student@cp:~$ cp <path_from_output_above>/CreateNFS.sh $HOME
```

```
student@cp:~$ bash $HOME/CreateNFS.sh
```

```
Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu xenial InRelease
Get:2 http://us-central1.gce.archive.ubuntu.com/ubuntu xenial-updates InRelease [102 kB]

<output_omitted>

Should be ready. Test here and second node

Export list for localhost:
/opt/sfw *
```

3. Test by mounting the resource from your **second node**. Begin by installing the client software.

```
student@worker:~$ sudo apt-get -y install nfs-common nfs-kernel-server
```

```
<output_omitted>
```

4. Test you can see the exported directory using **showmount** from you second node.

```
student@worker:~$ showmount -e cp #<-- Edit to be first node's name or IP
```

```
Export list for cp:
/opt/sfw *
```

5. Mount the directory. Be aware that unless you edit `/etc/fstab` this is not a persistent mount. Change out the node name for that of your cp node.

```
student@worker:~$ sudo mount cp:/opt/sfw /mnt
```

6. Verify the `hello.txt` file created by the script can be viewed.

```
student@worker:~$ ls -l /mnt
```

```
total 4
-rw-r--r-- 1 root root 9 Sep 28 17:55 hello.txt
```

7. Return to the cp node and create a YAML file for an object with kind **PersistentVolume**. The included example file needs an edit to the `server:` parameter. Use the hostname of the cp server and the directory you created in the previous step. Only syntax is checked, an incorrect name or directory will not generate an error, but a Pod using the incorrect resource will not start. Note that the `accessModes` do not currently affect actual access and are typically used as labels instead.

```
student@cp:~$ find $HOME -name PVol.yaml
```

```
<some_long_path>/PVol.yaml
```

```
student@cp:~$ cp <path_output_from_above>/PVol.yaml $HOME
```

```
student@cp:~$ vim PVol.yaml
```



PVol.yaml

```

1  apiVersion: v1
2  kind: PersistentVolume
3  metadata:
4    name: pvvol-1
5  spec:
6    capacity:
7      storage: 1Gi
8    accessModes:
9      - ReadWriteMany
10   persistentVolumeReclaimPolicy: Retain
11   nfs:
12     path: /opt/sfw
13     server: cp                                #<-- Edit to match cp node name or IP
14     readOnly: false

```

8. Create and verify you have a new 1Gi volume named **pvvol-1**. Note the status shows as Available. Remember we made two persistent volumes for the image registry earlier.

```
student@cp:~$ kubectl create -f PVol.yaml
```

```
persistentvolume/pvvol-1 created
```

```
student@cp:~$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS
↪ REASON AGE						
pvvol-1	1Gi	RWX	Retain	Available		
↪ 4s						
registryvm	200Mi	RWO	Retain	Bound	default/nginx-claim0	
↪ 4d						
task-pv-volume	200Mi	RWO	Retain	Bound	default/registry-claim0	
↪ 4d						

9. Now that we have a new volume we will use a **persistent volume claim (pvc)** to use it in a Pod. We should have two existing claims from our local registry.

```
student@cp:~/$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
nginx-claim0	Bound	registryvm	200Mi	RWO		4d
registry-claim0	Bound	task-pv-volume	200Mi	RWO		4d

10. Create or copy a yaml file with the kind **PersistentVolumeClaim**.

```
student@cp:~$ vim pvc.yaml
```



pvc.yaml

```

1  apiVersion: v1
2  kind: PersistentVolumeClaim
3  metadata:
4    name: pvc-one
5  spec:
6    accessModes:
7      - ReadWriteMany
8    resources:

```

YAML

```

9     requests:
10    storage: 200Mi

```

11. Create and verify the new pvc status is bound. Note the size is 1Gi, even though 200Mi was suggested. Only a volume of at least that size could be used, the first volume with found with at least that much space was chosen.

```
student@cp:~$ kubectl create -f pvc.yaml
```

```
persistentvolumeclaim/pvc-one created
```

```
student@cp:~$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
nginx-claim0	Bound	registryvm	200Mi	RWO		4d
pvc-one	Bound	pvvol-1	1Gi	RWX		4s
registry-claim0	Bound	task-pv-volume	200Mi	RWO		4d

12. Now look at the status of the physical volume. It should also show as bound.

```
student@cp:~$ kubectl get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS
CLAIM	STORAGECLASS	REASON	AGE	
pvvol-1	1Gi	RWX	Retain	Bound
default/pvc-one			14m	
registryvm	200Mi	RWO	Retain	Bound
default/nginx-claim0			4d	
task-pv-volume	200Mi	RWO	Retain	Bound
default/registry-claim0			4d	

13. Edit the `simpleapp.yaml` file to include two new sections. One section for the container while will use the volume mount point, you should have an existing entry for `car-vol`. The other section adds a volume to the deployment in general, which you can put after the `configMap` volume section.

```
student@cp:~$ vim $HOME/app1/simpleapp.yaml
```

YAML**simpleapp.yaml**

```

1  ....
2      volumeMounts:
3      - name: car-vol
4        mountPath: /etc/cars
5      - name: nfs-vol
6        mountPath: /opt
7  ....
8      volumes:
9      - name: car-vol
10        configMap:
11          defaultMode: 420
12          name: fast-car
13      - name: nfs-vol
14        persistentVolumeClaim:
15          claimName: pvc-one
16  status:
17  ....

```

#<-- Add this and following line

#<-- Add this and following two lines

14. Delete and re-create the deployment.

```
student@cp:~$ kubectl delete deployment try1 ; kubectl create -f $HOME/app1/simpleapp.yaml
```

```
deployment.apps "try1" deleted
deployment.apps/try1 created
```

15. View the details any of the pods in the deployment, you should see `nfs-vol` mounted under `/opt`. The use to command line completion with the `tab` key can be helpful for using a pod name.

```
student@cp:~$ kubectl describe pod try1-594fbb5fc7-5k7sj
```

```
<output_omitted>
Mounts:
  /etc/cars from car-vol (rw)
  /opt from nfs-vol (rw)
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-j7cqd (ro)
<output_omitted>
```

✍ Exercise 5.3: Using ConfigMaps Configure Containers

In an earlier lab we added a second container to handle logging. Now that we have learned about using ConfigMaps and attaching storage we will use configure our `basic` pod.

1. Review the YAML for our earlier simple pod. Recall that we added an Ambassador style logging container to the pod but had not fully configured the logging.

```
student@cp:~$ cat basic.yaml
```

```
<output_omitted>
containers:
- name: webcont
  image: nginx
  ports:
  - containerPort: 80
- name: fdlogger
  image: fluent/fluentd
```

2. Let us begin by adding shared storage to each container. We will use the `hostPath` storage class to provide the PV and PVC. First we create the directory.

```
student@cp:~$ sudo mkdir /tmp/weblog
```

3. Now we create a new PV to use that directory for the `hostPath` storage class. We will use the `storageClassName` of `manual` so that only PVCs which use that name will bind the resource.

```
student@cp:~$ vim weblog-pv.yaml
```

YAML

weblog-pv.yaml

```
1 kind: PersistentVolume
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-volume
5   labels:
6     type: local
7 spec:
```



```

8   storageClassName: manual
9   capacity:
10    storage: 100Mi
11   accessModes:
12    - ReadWriteOnce
13   hostPath:
14    path: "/tmp/weblog"

```

4. Create and verify the new PV exists and shows an Available status.

```
student@cp:~$ kubectl create -f weblog-pv.yaml
```

```
persistentvolume/weblog-pv-volume created
```

```
student@cp:~$ kubectl get pv weblog-pv-volume
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY
STATUS	CLAIM	STORAGECLASS	REASON AGE
weblog-pv-volume	100Mi	RWO	Retain
Available		manual	21s

5. Next we will create a PVC to use the PV we just created.

```
student@cp:~$ vim weblog-pvc.yaml
```



weblog-pvc.yaml

```

1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-claim
5 spec:
6   storageClassName: manual
7   accessModes:
8     - ReadWriteOnce
9   resources:
10    requests:
11     storage: 100Mi

```

6. Create the PVC and verify it shows as Bound to the the PV we previously created.

```
student@cp:~$ kubectl create -f weblog-pvc.yaml
```

```
persistentvolumeclaim/weblog-pv-claim created
```

```
student@cp:~$ kubectl get pvc weblog-pv-claim
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
weblog-pv-claim	Bound	weblog-pv-volume	100Mi	RWO
manual	79s			

7. We are ready to add the storage to our pod. We will edit three sections. The first will declare the storage to the pod in general, then two more sections which tell each container where to make the volume available.

```
student@cp:~$ vim basic.yaml
```

YAML
basic.yaml

```

1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: basicpod
5    labels:
6      type: webserver
7  spec:
8    volumes:                                #<-- Add three lines, same depth as containers
9      - name: weblog-pv-storage
10        persistentVolumeClaim:
11          claimName: weblog-pv-claim
12    containers:
13      - name: webcont
14        image: nginx
15        ports:
16          - containerPort: 80
17        volumeMounts:                      #<-- Add three lines, same depth as ports
18          - mountPath: "/var/log/nginx/"
19            name: weblog-pv-storage        # Must match volume name above
20      - name: fdlogger
21        image: fluent/fluentd
22        volumeMounts:                      #<-- Add three lines, same depth as image:
23          - mountPath: "/var/log"
24            name: weblog-pv-storage        # Must match volume name above

```

8. At this point we can create the pod again. When we create a shell we will find that the `access.log` for **nginx** is no longer a symbolic link pointing to `stdout` it is a writable, zero length file. Leave a **tailf** of the log file running.

```
student@cp:~$ kubectl create -f basic.yaml
```

```
pod/basicpod created
```

```
student@cp:~$ kubectl exec -c webcont -it basicpod -- /bin/bash
```



On Container

```
root@basicpod:/# ls -l /var/log/nginx/access.log
```

```
-rw-r--r-- 1 root root 0 Oct 18 16:12 /var/log/nginx/access.log
```

```
root@basicpod:/# tail -f /var/log/nginx/access.log
```

9. Open a second connection to your cp node. We will use the pod IP as we have not yet configured a service to expose the pod.

```
student@cp:~$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE						
basicpod	2/2	Running	0	3m26s	192.168.213.181	cp
<none>						

10. Use **curl** to view the welcome page of the webserver. When the command completes you should see a new entry added

to the log. Right after the GET we see a 200 response indicating success. You can use **ctrl-c** and **exit** to return to the host shell prompt.

```
student@cp:~$ curl http://192.168.213.181
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```



On Container

```
192.168.32.128 - - [18/Oct/2018:16:16:21 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.47.0" "-"
```

- Now that we know the webcont container is writing to the PV we will configure the logger to use that directory as a source. For greater flexibility we will configure **fluentd** using a configMap.

Fluentd has many options for input and output of data. We will read from a file of the webcont container and write to standard out of the fdlogger container. The details of the data settings can be found in **fluentd** documentation here: <https://docs.fluentd.org/v1.0/categories/config-file>

```
student@cp:~$ vim weblog-configmap.yaml
```



weblog-configmap.yaml

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: fluentd-config
5    namespace: default
6  data:
7    fluentd.conf: |
8      <source>
9        @type tail
10       format none
11       path /var/log/access.log
12       tag count.format1
13     </source>
14
15     <match *.**>
16       @type stdout
17       id stdout_output
18     </match>
```

- Create the new configMap.

```
student@cp:~$ kubectl create -f weblog-configmap.yaml
```

```
configmap/fluentd-config created
```

- View the logs for both containers in the basicpod. You should see some startup information, but not the HTTP traffic.

```
student@cp:~$ kubectl logs basicpod webcont
```

```
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
```

```
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
```

```
student@cp:~$ kubectl logs basicpod fdlogger
```

```
2020-09-02 19:32:59 +0000 [info]: reading config file path="/etc/fluentd-config/fluentd.conf"
2020-09-02 19:32:59 +0000 [info]: starting fluentd-0.12.29
2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-config-placeholders' version '0.4.0'
2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-plaintextformatter' version '0.2.6'

<output_omitted>

<source>
  @type tail
  format none
  path /var/log/access.log
<output_omitted>
```

14. Now we will edit the pod yaml file so that the **fluentd** container will mount the configmap as a volume and reference the variables inside the config file. You will add three areas, the volume declaration to the pod, the env parameter and the mounting of the volume to the fluentd container

```
student@cp:~$ vim basic.yaml
```

YAML

basic.yaml

```
1  ....
2  volumes:
3    - name: weblog-pv-storage
4      persistentVolumeClaim:
5        claimName: weblog-pv-claim
6    - name: log-config                                #<-- This and two lines following
7      configMap:
8        name: fluentd-config                          # Must match existing configMap
9  ....
10  image: fluent/fluentd
11  env:                                                #<-- This and two lines following
12    - name: FLUENTD_OPT
13      value: -c /etc/fluentd-config/fluentd.conf
14  ....
15  volumeMounts:
16    - mountPath: "/var/log"
17      name: weblog-pv-storage
18    - name: log-config                                #<-- This and next line
19      mountPath: "/etc/fluentd-config"
```

15. At this point we can delete and re-create the pod, which would cause the configmap to be used by the new pod, among other changes.

```
student@cp:~$ kubectl delete pod basicpod
```

```
pod "basicpod" deleted
```

```
student@cp:~$ kubectl create -f basic.yaml
```

```
pod/basicpod created
```

```
student@cp:~$ kubectl get pod basicpod -o wide
```


NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED....
basicpod	2/2	Running	0	8s	192.168.171.122	worker	<none>

16. Use **curl** a few times to look at the default page served by basicpod

```
student@cp:~$ curl http://192.168.171.122
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
<output_omitted>
```

17. Look at the logs for both containers. In addition to the standard startup information, you should also see the HTTP requests from the curl commands you just used at the end of the fdlogger output.

```
student@cp:~$ kubectl logs basicpod webcont
```

```
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
```

```
student@cp:~$ kubectl logs basicpod fdlogger
```

```
2020-09-02 19:32:59 +0000 [info]: reading config file path="/etc/fluentd-config/fluentd.conf"
2020-09-02 19:32:59 +0000 [info]: starting fluentd-0.12.29
2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-config-placeholders' version '0.4.0'
2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-plaintextformatter' version '0.2.6'

<output_omitted>

<source>
  @type tail
  format none
  path /var/log/access.log

<output_omitted>

2020-09-02 19:47:38 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:38
↪ +0000] \"GET / HTTP/1.1\" 200 612 \"-\" \"curl/7.58.0\" \"-\""}
2020-09-02 19:47:41 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:41
↪ +0000] \"GET / HTTP/1.1\" 200 612 \"-\" \"curl/7.58.0\" \"-\""}
2020-09-02 19:47:47 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:47
↪ +0000] \"GET / HTTP/1.1\" 200 612 \"-\" \"curl/7.58.0\" \"-\""}

```

Exercise 5.4: Rolling Updates and Rollbacks

We created our simpleapp image using **podman**, and will create an update to the container. Remember that if you are using **docker** the command syntax should be the same.

1. Make a slight change to our source and create a new image. We will use updates and rollbacks with our application. Adding a comment to the last line should be enough for a new image to be generated.

```
student@cp:~$ cd $HOME/app1
```

```
student@cp:~/app1$ vim simple.py
```

```
<output_omitted>
```

```
## Sleep for five seconds then continue the loop
```

```
time.sleep(5)
```

```
## Adding a new comment so image is different.
```

2. View the current images, tags and how long ago the images were created.

```
student@cp:~/app1$ sudo podman images |grep simple
```

10.97.177.111:5000/simpleapp	latest	fd6e141c3a2e	14 hours ago	925 MB
localhost/simpleapp	latest	fd6e141c3a2e	14 hours ago	925 MB

3. Build the image again. A new container and image will be created. Verify when successful. There should be a different image ID and a recent creation time for the local image.

```
student@cp:~/app1$ sudo podman build -t simpleapp .
```

```
STEP 1: FROM python:2
STEP 2: ADD simple.py /
--> c1620c253cb
STEP 3: CMD [ "python", "./simple.py" ]
STEP 4: COMMIT simpleapp
--> fd46f4115eb
fd46f4115ebc87a064c15c57b08779b02deec97ab1abe92713689342ece496bd
```

```
student@cp:~/app1$ sudo podman images | grep simple
```

localhost/simpleapp	latest	fd46f4115ebc	43 seconds ago	925 MB
10.97.177.111:5000/simpleapp	latest	fd6e141c3a2e	14 hours ago	925 MB

4. Tag and push the updated image to your locally hosted registry. Use the tag v2 this time at the end of the image name.

```
student@cp:~/app1$ sudo podman tag simpleapp $repo/simpleapp:v2
```

```
student@cp:~/app1$ sudo podman push $repo/simpleapp:v2
```

```
Getting image source signatures
Copying blob 46829331b1e4 skipped: already exists
Copying blob 47458fb45d99 skipped: already exists
Copying blob 461719022993 skipped: already exists
Copying blob a3c1026c6bcc skipped: already exists
Copying blob d35c5bda4793 skipped: already exists
Copying blob f1d420c2af1a skipped: already exists
Copying blob bbcb11d3fa81 done
Copying blob ceee8816bb96 skipped: already exists
Copying blob da7b0a80a4f2 skipped: already exists
Copying blob e571d2d3c73c skipped: already exists
Copying config fd46f4115e done
Writing manifest to image destination
Storing signatures
```

5. Check your images again, there should be an additional entry with same image ID, but with the tag of v2.

```
student@cp:~/app1$ sudo podman images |grep simple
```

10.97.177.111:5000/simpleapp	v2	fd46f4115ebc	3 minutes ago	925 MB
localhost/simpleapp	latest	fd46f4115ebc	3 minutes ago	925 MB
10.97.177.111:5000/simpleapp	latest	fd6e141c3a2e	14 hours ago	925 MB

6. Connect to a terminal running on your worker node. Pull the image without asking for a version, which one would expect to pull the latest image, then pull v2. Note the default did not pull the new version of the image. Literally the latest tag has no relationship to the latest image.

```
student@worker:~$ sudo podman pull $repo/simpleapp
```

```
Trying to pull 10.97.177.111:5000/simpleapp:latest...
Getting image source signatures
Copying blob ca488b1eb9fa skipped: already exists
Copying blob 1f22e54987ac skipped: already exists
Copying blob 33be93024b52 skipped: already exists
Copying blob 41b428188c3d skipped: already exists
Copying blob 6dae3f0239bb skipped: already exists
Copying blob 637977187da0 skipped: already exists
Copying blob 3263dd1bdf84 [-----] 0.0b / 0.0b
Copying blob 7506b856c5df [-----] 0.0b / 0.0b
Copying blob 19aa3a14ed80 [-----] 0.0b / 0.0b
Copying blob c45f0d38c649 [-----] 0.0b / 0.0b
Copying config fd6e141c3a done
Writing manifest to image destination
Storing signatures
fd6e141c3a2e2d663f56c0fa7a4d056f73f6d1f63da6e36b19d7ef006dd8fff6
```

```
student@worker:~$ sudo podman pull $repo/simpleapp:v2
```

```
Trying to pull 10.97.177.111:5000/simpleapp:v2...
Getting image source signatures
Copying blob ca488b1eb9fa skipped: already exists
Copying blob 1f22e54987ac skipped: already exists
Copying blob 6dae3f0239bb skipped: already exists
Copying blob 41b428188c3d skipped: already exists
Copying blob 637977187da0 skipped: already exists
Copying blob 33be93024b52 skipped: already exists
Copying blob 3263dd1bdf84 skipped: already exists
Copying blob 7506b856c5df skipped: already exists
Copying blob c45f0d38c649 skipped: already exists
Copying blob 4da1215fbbb8 done
Copying config fd46f4115e done
Writing manifest to image destination
Storing signatures
fd46f4115ebc87a064c15c57b08779b02deec97ab1abe92713689342ece496bd
```

7. Return to your cp node, use **kubectl edit** to update the image for the try1 deployment to use v2. As we are only changing one parameter we could also use the **kubectl set** command. Note that the configuration file has not been updated, so a delete or a replace command would not include the new version. It can take the pods up to a minute to delete and to recreate each pod in sequence.

```
student@cp:~/app1$ kubectl edit deployment try1
```

```
....
containers:
- image: 10.105.119.236:5000/simpleapp:v2  #<-- Edit tag
  imagePullPolicy: Always
....
```

8. Verify each of the pods has been recreated and is using the new version of the image. Note some messages will show the scaling down of the old **replicaset**, others should show the scaling up using the new image.

```
student@cp:~/app1$ kubectl get events
```

```
42m      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fdbb5d557 to 6
32s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 2
32s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to
↪ 5
32s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 3
23s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to
↪ 4
23s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 4
22s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to
↪ 3
22s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 5
18s      Normal    ScalingReplicaSet   Deployment   Scaled down replica set try1-7fdbb5d557 to
↪ 2
18s      Normal    ScalingReplicaSet   Deployment   Scaled up replica set try1-7fd7459fc6 to 6
8s      Normal    ScalingReplicaSet   Deployment   (combined from similar events):
Scaled down replica set try1-7fdbb5d557 to 0
```

9. View the images of a Pod in the deployment. Narrow the output to just view the images. The goproxy remains unchanged, but the simpleapp should now be v2.

```
student@cp:~/app1$ kubectl describe pod try1-895fccfb-ttqdn |grep Image
```

```
Image:          10.105.119.236:5000/simpleapp:v2
Image ID:        10.97.177.111:5000/simpleapp@sha256:c3848fd3b5cdf5f241ceb4aa7e96c5c2be9f09
f23eealafe946f8507f3fcbc29
Image:          registry.k8s.io/goproxy:0.1
Image ID:        docker.io/library/goproxy@sha256:5334c7ad43048e3538775c
b09aaf184f5e8c
c8f1d93c209865c8f1d93c2098
```

10. View the update history of the deployment.

```
student@cp:~/app1$ kubectl rollout history deployment try1
```

```
deployment.apps/try1
REVISION  CHANGE-CAUSE
1          <none>
2          <none>
```

11. Compare the output of the **rollout history** for the two revisions. Images and labels should be different, with the image v2 being the change we made.

```
student@cp:~/app1$ kubectl rollout history deployment try1 --revision=1 > one.out
```

```
student@cp:~/app1$ kubectl rollout history deployment try1 --revision=2 > two.out
```

```
student@cp:~/app1$ diff one.out two.out
```

```
1c1
< deployment.apps/try1 with revision #1
---
> deployment.apps/try1 with revision #2
4c4
<         pod-template-hash=1509661973
---
>         pod-template-hash=45197796
7c7
<   Image:          10.105.119.236:5000/simpleapp
---
>   Image:          10.105.119.236:5000/simpleapp:v2
```

12. View what would be undone using the **dry-run** option while undoing the rollout. This allows us to see the new template prior to using it.

```
student@cp:~/app1$ kubectl rollout undo --dry-run=client deployment/try1
```

```
deployment.apps/try1 Pod Template:
  Labels:      app=try1
              pod-template-hash=754bf9c75b
  Containers:
    simpleapp:
      Image:    10.97.177.111:5000/simpleapp
      Port:     <none>
      Host Port: <none>
  <output_omitted>
```

13. View the pods. Depending on how fast you type the try1 pods should be about 2 minutes old.

```
student@cp:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	5d
registry-795c6c8b8f-hl5wf	1/1	Running	2	5d
try1-594fbb5fc7-7dl7c	2/2	Running	0	2m
try1-594fbb5fc7-8mxlb	2/2	Running	0	2m
try1-594fbb5fc7-jr7h7	2/2	Running	0	2m
try1-594fbb5fc7-s24wt	2/2	Running	0	2m
try1-594fbb5fc7-xfffg	2/2	Running	0	2m
try1-594fbb5fc7-zfmz8	2/2	Running	0	2m

14. In our case there are only two revisions, which is also the default number kept. Were there more we could choose a particular version. The following command would have the same effect as the previous, without the **dry-run** option.

```
student@cp:~/app1$ kubectl rollout undo deployment try1 --to-revision=1
```

```
deployment.apps/try1 rolled back
```

15. Again, it can take a bit for the pods to be terminated and re-created. Keep checking back until they are all running again.

```
student@cp:~/app1$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	5d
registry-795c6c8b8f-hl5wf	1/1	Running	2	5d
try1-594fbb5fc7-7dl7c	2/2	Terminating	0	3m
try1-594fbb5fc7-8mxlb	0/2	Terminating	0	2m
try1-594fbb5fc7-jr7h7	2/2	Terminating	0	3m
try1-594fbb5fc7-s24wt	2/2	Terminating	0	2m
try1-594fbb5fc7-xfffg	2/2	Terminating	0	3m
try1-594fbb5fc7-zfmz8	1/2	Terminating	0	2m
try1-895fccfb-8dn4b	2/2	Running	0	22s
try1-895fccfb-kz72j	2/2	Running	0	10s
try1-895fccfb-rxxtw	2/2	Running	0	24s
try1-895fccfb-srwq4	1/2	Running	0	11s
try1-895fccfb-vkvmb	2/2	Running	0	31s
try1-895fccfb-z46qr	2/2	Running	0	31s

Exercise 5.5: Domain Review

**Very Important**

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#). locate some of the topics we have covered in this chapter, which should match up with the following list of bullet points.

- Utilize persistent and ephemeral volumes
- Understand Deployments and how to perform rolling updates
- Understand ConfigMaps

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

Using only the allowed browser, URLs, and subdomains search for and bookmark a YAML example to create and configure the resources called for in this review.

1. Create a new secret called `specialofday` using the key `entree` and the value `meatloaf`.
2. Create a new deployment called `foodie` running the `nginx` image.
3. Add the `specialofday` secret to pod mounted as a volume under the `/food/` directory.
4. Execute a bash shell inside a `foodie` pod and verify the secret has been properly mounted.
5. Update the deployment to use the `nginx:1.12.1-alpine` image and verify the new image is in use.
6. Roll back the deployment and verify the typical, current stable version of `nginx` is in use again.
7. Create a new 200M NFS volume called `reviewvol` using the NFS server configured earlier in the lab.
8. Create a new PVC called `reviewpvc` which will uses the `reviewvol` volume.
9. Edit the deployment to use the PVC and mount the volume under `/newvol`
10. Execute a bash shell into the `nginx` container and verify the volume has been mounted.
11. Delete any resources created during this review.

Chapter 6

Understanding Security



Exercise 6.1: Set SecurityContext for a Pod and Container

Working with Security: Overview

In this lab we will implement security features for new applications, as the simpleapp YAML file is getting long and more difficult to read. Kubernetes architecture favors smaller, decoupled, and transient applications working together. We'll continue to emulate that in our exercises.

In this exercise we will create two new applications. One will be limited in its access to the host node, but have access to encoded data. The second will use a `network security policy` to move from the default all-access Kubernetes policies to a mostly closed network. First we will set `security contexts` for pods and containers, then create and consume secrets, then finish with configuring a network security policy.

1. Begin by making a new directory for our second application. Change into that directory.

```
student@cp:~$ mkdir $HOME/app2
```

```
student@cp:~$ cd $HOME/app2/
```

2. Create a YAML file for the second application. In the example below we are using a simple image, `busybox`, which allows access to a shell, but not much more. We will add a `runAsUser` to both the pod as well as the container.

```
student@cp:~/app2$ vim second.yaml
```

YAML

second.yaml

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: secondapp
5  spec:
6    securityContext:
7      runAsUser: 1000
8    containers:
9      - name: busy
10      image: busybox
```



```

11     command:
12         - sleep
13         - "3600"
14     securityContext:
15         runAsUser: 2000
16         allowPrivilegeEscalation: false

```

3. Create the secondapp pod and verify it's running. Unlike the previous deployment this application is running as a pod. Look at the YAML output, to compare and contrast with what a deployment looks like. The status section probably has the largest contrast.

```
student@cp:~/app2$ kubectl create -f second.yaml
```

```
pod/secondapp created
```

```
student@cp:~/app2$ kubectl get pod secondapp
```

NAME	READY	STATUS	RESTARTS	AGE
secondapp	1/1	Running	0	21s

```
student@cp:~/app2$ kubectl get pod secondapp -o yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  annotations:
    cni.projectcalico.org/containerID: f706f1c343de8749cb687db40....
    cni.projectcalico.org/podIP: 192.168.158.97/32
    cni.projectcalico.org/podIPs: 192.168.158.97/32
  creationTimestamp: "2021-05-03T21:23:12Z"
  name: secondapp
  <output_omitted>

```

4. Execute a Bourne shell within the Pod. Check the user ID of the shell and other processes. It should show the container setting, not the pod. This allows for multiple containers within a pod to customize their UID if desired. As there is only one container in the pod we do not need to use the **-c busy** option.

```
student@cp:~/app2$ kubectl exec -it secondapp -- sh
```



On Container

```
/ $ ps aux
```

PID	USER	TIME	COMMAND
1	2000	0:00	sleep 3600
8	2000	0:00	sh
12	2000	0:00	ps aux

5. While here check the capabilities of the kernel. In upcoming steps we will modify these values.



On Container

```
/ $ grep Cap /proc/1/status
```




```
CapInh:      000000000000005fb
CapPrm:      0000000000000000
CapEff:      0000000000000000
CapBnd:      000000000000005fb
CapAmb:      0000000000000000
```

```
/ $ exit
```

6. Use the capability shell wrapper tool, the **capsh** command, to decode the output. We will view and compare the output in a few steps. Note that there are something like nine comma separated capabilities listed. The number may change in future versions.

```
student@cp:~/app2$ capsh --decode=000000000000005fb
```

```
0x000000000000005fb=cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,
cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service
```

7. Edit the YAML file to include new capabilities for the container. A capability allows granting of specific, elevated privileges without granting full root access. We will be setting **NET_ADMIN** to allow interface, routing, and other network configuration. We'll also set **SYS.TIME**, which allows system clock configuration. More on kernel capabilities can be read here: <https://github.com/torvalds/linux/blob/master/include/uapi/linux/capability.h>

It can take up to a minute for the pod to fully terminate, allowing the future pod to be created.

```
student@cp:~/app2$ kubectl delete pod secondapp
```

```
pod "secondapp" deleted
```

```
student@cp:~/app2$ vim second.yaml
```

YAML

second.yaml

```
1 <output_omitted>
2   - sleep
3   - "3600"
4   securityContext:
5     runAsUser: 2000
6     allowPrivilegeEscalation: false
7     capabilities:                                #<-- Add this and following line
8     add: ["NET_ADMIN", "SYS_TIME"]
```

8. Create the pod again. Execute a shell within the container and review the Cap settings under `/proc/1/status`. They should be different from the previous instance.

```
student@cp:~/app2$ kubectl create -f second.yaml
```

```
pod/secondapp created
```

```
student@cp:~/app2$ kubectl exec -it secondapp -- sh
```



On Container

```
/ $ grep Cap /proc/1/status
```



```
CapInh:      00000000020015fb
CapPrm:      0000000000000000
CapEff:      0000000000000000
CapBnd:      00000000020015fb
CapAmb:      0000000000000000
```

```
/ $ exit
```

9. Decode the output again. Note that the instance now has 16 comma delimited capabilities listed. **cap_net_admin** is listed as well as **cap_sys_time**.

```
student@cp:~/app2$ capsh --decode=00000000020015fb
```

```
0x00000000020015fb=cap_chown,cap_dac_override,cap_fowner,cap_fsetid,cap_kill,
cap_setgid,cap_setuid,cap_setpcap,cap_net_bind_service,cap_net_admin,cap_sys_time
```

✍ Exercise 6.2: Create and consume Secrets

Secrets are consumed in a manner similar to ConfigMaps, covered in an earlier lab. While at-rest encryption is now easy to configure, historically a secret was just base64 encoded. There are several types of encryption which can be configured.

1. Begin by generating an encoded password.

```
student@cp:~/app2$ echo LFTr@1n | base64
```

```
TEZUckAxbgo=
```

2. Create a YAML file for the object with an API object kind set to Secret. Use the encoded key as a password parameter.

```
student@cp:~/app2$ vim secret.yaml
```

YAML

secret.yaml

```
1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: lfsecret
5 data:
6   password: TEZUckAxbgo=
```

3. Ingest the new object into the cluster.

```
student@cp:~/app2$ kubectl create -f secret.yaml
```

```
secret/lfsecret created
```

4. Edit secondapp YAML file to use the secret as a volume mounted under `/mysqlpassword`. `volumeMounts:` lines up with the container name: and `volumes:` lines up with containers: Note the pod will restart when the sleep command finishes every 3600 seconds, or every hour.

```
student@cp:~/app2$ vim second.yaml
```



second.yaml

```

1 ....
2     runAsUser: 2000
3     allowPrivilegeEscalation: false
4     capabilities:
5         add: ["NET_ADMIN", "SYS_TIME"]
6     volumeMounts:                                #<-- Add this and six following lines
7     - name: mysql
8       mountPath: /mysqlpassword
9     volumes:
10    - name: mysql
11      secret:
12        secretName: lfsecret

```

```
student@cp:~/app2$ kubectl delete pod secondapp
```

```
pod "secondapp" deleted
```

```
student@cp:~/app2$ kubectl create -f second.yaml
```

```
pod/secondapp created
```

5. Verify the pod is running, then check if the password is mounted where expected. We will find that the password is available in its clear-text, decoded state.

```
student@cp:~/app2$ kubectl get pod secondapp
```

NAME	READY	STATUS	RESTARTS	AGE
secondapp	1/1	Running	0	34s

```
student@cp:~/app2$ kubectl exec -ti secondapp -- /bin/sh
```



On Container

```
/ $ cat /mysqlpassword/password
```

```
LFTTr@1n
```

6. View the location of the directory. Note it is a symbolic link to `..data` which is also a symbolic link to another directory. After taking a look at the filesystem within the container, exit back to the node.



On Container

```
/ $ cd /mysqlpassword/
```

```
/mysqlpassword $ ls
```

```
password
```

```
/mysqlpassword $ ls -al
```



```
total 4
drwxrwxrwt  3 root    root      100 May 19 16:06 .
dr-xr-xr-x  1 root    root      4096 May 19 16:06 ..
drwxr-xr-x  2 root    root        60 May 19 16:06 ..2021_05_19_16_06_41.694089911
lrwxrwxrwx  1 root    root        31 May 19 16:06 ..data ->
-> ..2021_05_19_16_06_41.694089911
lrwxrwxrwx  1 root    root        15 May 19 16:06 password -> ..data/password
```

```
/mysqlpassword $ exit
```

✍ Exercise 6.3: Working with ServiceAccounts

We can use ServiceAccounts to assign cluster roles, or the ability to use particular HTTP verbs. In this section we will create a new ServiceAccount and grant it access to view secrets.

1. Begin by viewing secrets, both in the default namespace as well as all.

```
student@cp:~/app2$ cd
```

```
student@cp:~$ kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-c4rdg	kubernetes.io/service-account-token	3	4d16h
lfsecret	Opaque	1	6m5s

```
student@cp:~$ kubectl get secrets --all-namespaces
```

NAMESPACE	NAME	TYPE	DATA	AGE
default	default-token-c4rdg	kubernetes.io/service-account-token	3	4d16h
kube-public	default-token-zqzbg	kubernetes.io/service-account-token	3	4d16h
kube-system	attachdetach-controller-token-wxzvc	kubernetes.io/service-account-token	3	4d16h

<output_omitted>

2. We can see that each agent uses a secret in order to interact with the API server. We will create a new ServiceAccount which will have access.

```
student@cp:~$ vim serviceaccount.yaml
```



serviceaccount.yaml

```
1 apiVersion: v1
2 kind: ServiceAccount
3 metadata:
4   name: secret-access-sa
```

```
student@cp:~$ kubectl create -f serviceaccount.yaml
```

```
serviceaccount/secret-access-sa created
```

```
student@cp:~$ kubectl get serviceaccounts
```

NAME	SECRETS	AGE
default	1	1d17h
secret-access-sa	1	34s

3. Now we will create a ClusterRole which will list the actual actions allowed cluster-wide. We will look at an existing role to see the syntax.

```
student@cp:~$ kubectl get clusterroles
```

NAME	AGE
admin	1d17h
calico-cni-plugin	1d17h
calico-kube-controllers	1d17h
cluster-admin	1d17h
<output_omitted>	

4. View the details for the admin and compare it to the cluster-admin. The admin has particular actions allowed, but cluster-admin has the meta-character '*' allowing all actions.

```
student@cp:~$ kubectl get clusterroles admin -o yaml
```

```
<output_omitted>
```

```
student@cp:~$ kubectl get clusterroles cluster-admin -o yaml
```

```
<output_omitted>
```

5. Using some of the output above, we will create our own file.

```
student@cp:~$ vim clusterrole.yaml
```

YAML

clusterrole.yaml

```
1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: ClusterRole
3 metadata:
4   name: secret-access-cr
5 rules:
6 - apiGroups:
7   - ""
8   resources:
9     - secrets
10  verbs:
11    - get
12    - list
```

6. Create and verify the new ClusterRole.

```
student@cp:~$ kubectl create -f clusterrole.yaml
```

```
clusterrole.rbac.authorization.k8s.io/secret-access-cr created
```

```
student@cp:~$ kubectl get clusterrole secret-access-cr -o yaml
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
```

```
creationTimestamp: 2018-10-18T19:27:24Z
name: secret-access-cr
<output_omitted>
```

7. Now we bind the role to the account. Create another YAML file which uses `roleRef::`

```
student@cp:~$ vim rolebinding.yaml
```

YAML

rolebinding.yaml

```
1 apiVersion: rbac.authorization.k8s.io/v1
2 kind: RoleBinding
3 metadata:
4   name: secret-rb
5 subjects:
6 - kind: ServiceAccount
7   name: secret-access-sa
8 roleRef:
9   kind: ClusterRole
10  name: secret-access-cr
11 apiGroup: rbac.authorization.k8s.io
```

8. Create the new RoleBinding and verify.

```
student@cp:~$ kubectl create -f rolebinding.yaml
```

```
rolebinding.rbac.authorization.k8s.io/secret-rb created
```

```
student@cp:~$ kubectl get rolebindings
```

```
NAME      AGE
secret-rb 17s
```

9. View the `secondapp` pod and **grep** for the current `serviceAccount`. Note that it uses the default account.

```
student@cp:~$ kubectl get pod secondapp -o yaml |grep serviceAccount
```

```
serviceAccount: default
serviceAccountName: default
- serviceAccountToken:
```

10. Edit the `second.yaml` file and add the use of the `serviceAccount`.

```
student@cp:~$ vim $HOME/app2/second.yaml
```

YAML

second.yaml

```
1 ....
2   name: secondapp
3 spec:
4   serviceAccountName: secret-access-sa #<-- Add this line
5   securityContext:
6     runAsUser: 1000
7   ....
```

11. We will delete the `secondapp` pod if still running, then create it again. Note that the `serviceAccount` is no longer the default.

```
student@cp:~$ kubectl delete pod secondapp ; kubectl create -f $HOME/app2/second.yaml
```

```
pod "secondapp" deleted
pod/secondapp created
```

```
student@cp:~$ kubectl get pod secondapp -o yaml | grep serviceAccount
```

```
serviceAccount: secret-access-sa
serviceAccountName: secret-access-sa
- serviceAccountToken:
```

Exercise 6.4: Implement a NetworkPolicy

An early architecture decision with Kubernetes was non-isolation, that all pods were able to connect to all other pods and nodes by design. In more recent releases the use of a `NetworkPolicy` allows for pod isolation. The policy only has effect when the network plugin, like **Project Calico**, are capable of honoring them. If used with a plugin like **flannel** they will have no effect. The use of `matchLabels` allows for more granular selection within the namespace which can be selected using a `namespaceSelector`. Using multiple labels can allow for complex application of rules. More information can be found here: <https://kubernetes.io/docs/concepts/services-networking/network-policies>

1. Begin by creating a default policy which denies all traffic. Once ingested into the cluster this will affect every pod not selected by another policy, creating a mostly-closed environment. If you want to only deny ingress or egress traffic you can remove the other `policyType`.

```
student@cp:~$ cd $HOME/app2/
```

```
student@cp:~/app2$ vim allclosed.yaml
```

YAML

allclosed.yaml

```
1 apiVersion: networking.k8s.io/v1
2 kind: NetworkPolicy
3 metadata:
4   name: deny-default
5 spec:
6   podSelector: {}
7   policyTypes:
8   - Ingress
9   - Egress
```

2. Before we can test the new network policy we need to make sure network access works without it applied. Update **secondapp** to include a new container running **nginx**, then test access. Begin by adding two lines for the **nginx** image and name `webserver`, as found below. It takes a bit for the pod to terminate, so we'll delete then edit the file.

```
student@cp:~/app2$ kubectl delete pod secondapp
```

```
pod "secondapp" deleted
```

```
student@cp:~/app2$ vim second.yaml
```



second.yaml

```

1 spec:
2   serviceAccountName: secret-access-sa
3   securityContext:
4     runAsUser: 1000
5   containers:
6     - name: webserver                                #<-- Add this and following line
7       image: nginx
8     - name: busy
9       image: busybox
10      command:

```

3. Create the new pod. Be aware the pod will move from ContainerCreating to Error to CrashLoopBackOff, as only one of the containers will start. We will troubleshoot the error in following steps.

```
student@cp:~/app2$ kubectl create -f second.yaml
```

```
pod/secondapp created
```

```
student@cp:~/app2$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-6b58d9cdfd-9fnl4	1/1	Running	1	2d
Registry-795c6c8b8f-hl5wf	1/1	Running	2	2d
secondapp	1/2	CrashLoopBackOff	1	13s

<output_omitted>

4. Take a closer look at the events leading up to the failure. The images were pulled and the container was started. It was the full execution of the container which failed.

```
student@cp:~/app2$ kubectl get event
```

```

<output_omitted>
25s      Normal    Scheduled    Pod    Successfully assigned default/secondapp to cp
4s       Normal    Pulling     Pod    pulling image "nginx"
2s       Normal    Pulled      Pod    Successfully pulled image "nginx"
2s       Normal    Created     Pod    Created container
2s       Normal    Started     Pod    Started container
23s      Normal    Pulling     Pod    pulling image "busybox"
21s      Normal    Pulled      Pod    Successfully pulled image "busybox"
21s      Normal    Created     Pod    Created container
21s      Normal    Started     Pod    Started container
1s       Warning   BackOff     Pod    Back-off restarting failed container

```

5. View the logs of the **webserver** container mentioned in the previous output. Note there are errors about the user directive and not having permission to make directories.

```
student@cp:~/app2$ kubectl logs secondapp webserver
```

```

2018/04/13 19:51:13 warn 1 the "user" directive makes sense
only if the cp process runs with super-user privileges,
  ignored in /etc/nginx/nginx.conf:2
nginx: warn the "user" directive makes sense only if the cp
process runs with super-user privileges,
  ignored in /etc/nginx/nginx.conf:2
2018/04/13 19:51:13 emerg 11: mkdir(\\) "/var/cache/nginx/client temp"
failed (13: Permission denied)
nginx: emerg mkdir() "/var/cache/nginx/client\_temp" failed
(13: Permission denied)

```


6. Delete the pods. Edit the YAML file to comment out the setting of a UID for the entire pod.

```
student@cp:~/app2$ kubectl delete -f second.yaml
```

```
pod "secondapp" deleted
```

```
student@cp:~/app2$ vim second.yaml
```

YAML

```
1 spec:
2   serviceAccountName: secret-access-sa
3   # securityContext:                                #<-- Comment this and following line
4   #   runAsUser: 1000
5   containers:
6     - name: webserver
```

7. Create the pod again. This time both containers should run. You may have to wait for the previous pod to fully terminate, depending on how fast you type.

```
student@cp:~/app2$ kubectl create -f second.yaml
```

```
pod/secondapp created
```

```
student@cp:~/app2$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
<output_omitted>				
secondapp	2/2	Running	0	5s

8. Expose the **webserver** using a NodePort service. Expect an error due to lack of labels.

```
student@cp:~/app2$ kubectl expose pod secondapp --type=NodePort --port=80
```

```
error: couldn't retrieve selectors via --selector flag or
introspection: the pod has no labels and cannot be exposed
See 'kubectl expose -h' for help and examples.
```

9. Edit the YAML file to add a label in the metadata, adding the example: second label right after the pod name. Note you can delete several resources at once by passing the YAML file to the delete command. Delete and recreate the pod. It may take up to a minute for the pod to shut down.

```
student@cp:~/app2$ kubectl delete -f second.yaml
```

```
pod "secondapp" deleted
```

```
student@cp:~/app2$ vim second.yaml
```

YAML

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: secondapp
5   labels:                                #<-- This and following line
6     example: second
7 spec:
8   # securityContext:
9   #   runAsUser: 1000
10  <output_omitted>
```

```
student@cp:~/app2$ kubectl create -f second.yaml
```

```
pod/secondapp created
```

```
student@cp:~/app2$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
<output_omitted>				
secondapp	2/2	Running	0	15s

10. This time we will expose a NodePort again, and create the service separately, then add a label to illustrate how labels are essential for tying resources together inside of kubernetes.

```
student@cp:~/app2$ kubectl create service nodeport secondapp --tcp=80
```

```
service/secondapp created
```

11. Look at the details of the service. Note the selector is set to app: secondapp. Also take note of the nodePort, which is 31655 in the example below, yours may be different.

```
student@cp:~/app2$ kubectl get svc secondapp -o yaml
```

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: "2020-04-16T04:40:07Z"
  labels:
    example: second
  managedFields:
  ....
spec:
  clusterIP: 10.97.96.75
  externalTrafficPolicy: Cluster
  ports:
  - nodePort: 31655
    port: 80
    protocol: TCP
    targetPort: 80
  selector:
    app: secondapp
  sessionAffinity: None
  type: NodePort
status:
  loadBalancer: {}
```

12. Test access to the service using **curl** and the ClusterIP shown in the previous output. As the label does not match any other resources, the **curl** command should fail. If it hangs **control-c** to exit back to the shell.

```
student@cp:~/app2$ curl http://10.97.96.75
```

13. Edit the service. We will change the label to match **secondapp**, and set the nodePort to a new port, one that may have been specifically opened by our firewall team, port 32000.

```
student@cp:~/app2$ kubectl edit svc secondapp
```

YAML

```
1 <output_omitted>
2 ports:
```



```

3 - name: "80"
4   nodePort: 32000      #<-- Edit this line
5   port: 80
6   protocol: TCP
7   targetPort: 80
8   selector:
9     example: second    #<-- Edit this line, note key and parameter change
10  sessionAffinity: None
11 <output_omitted>

```

14. Verify the updated port number is showing properly, and take note of the ClusterIP. The example below shows a ClusterIP of 10.97.96.75 and a port of 32000 as expected.

```
student@cp:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
<output_omitted>					
secondapp	NodePort	10.97.96.75	<none>	80:32000/TCP	5m

15. Test access to the high port. You should get the default nginx welcome page both if you test from the node to the ClusterIP:<low-port-number> and from the exterior hostIP:<high-port-number>. As the high port is randomly generated make sure it's available. Both of your nodes should be exposing the web server on port 32000. The example shows the use of the **curl** command, you could also use a web browser.

```
student@cp:~/app2$ curl http://10.97.96.75
```

```

<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```

```
student@cp:~/app2$ curl ifconfig.io
```

```
35.184.219.5
```

```
[user@laptop ~]$ curl http://35.184.219.5:32000
```

```

<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```

16. Now test egress from a container to the outside world. We'll use the **netcat** command to verify access to a running web server on port 80. First test local access to nginx, then a remote server.

```
student@cp:~/app2$ kubectl exec -it -c busy secondapp -- sh
```



On Container

```
/ $ nc -vz 127.0.0.1 80
```

```
127.0.0.1 (127.0.0.1:80) open
```

```
/ $ nc -vz www.linux.com 80
```



```
www.linux.com (151.101.185.5:80) open
```

```
/ $ exit
```

✍ Exercise 6.5: Testing the Policy

1. Now that we have tested both ingress and egress we can implement the network policy.

```
student@cp:~/app2$ kubectl create -f $HOME/app2/allclosed.yaml
```

```
networkpolicy.networking.k8s.io/deny-default created
```

2. Use the ingress and egress tests again. Three of the four should eventually timeout. Start by testing from outside the cluster, and interrupt if you get tired of waiting.

```
[user@laptop ~]$ curl http://35.184.219.5:32000
```

```
curl: (7) Failed to connect to 35.184.219.5 port
32000: Connection timed out
```

3. Then test from the host to the container.

```
student@cp:~/app2$ curl http://10.97.96.75:80
```

```
curl: (7) Failed to connect to 10.97.96.75 port 80: Connection timed out
```

4. Now test egress. From container to container should work, as the filter is outside of the pod. Then test egress to an external web page. It should eventually timeout.

```
student@cp:~/app2$ kubectl exec -it -c busy secondapp -- sh
```



On Container

```
/ $ nc -vz 127.0.0.1 80
```

```
127.0.0.1 (127.0.0.1:80) open
```

```
/ $ nc -vz www.linux.com 80
```

```
nc: bad address 'www.linux.com'
```

```
/ $ exit
```

5. Update the NetworkPolicy and comment out the Egress line. Then replace the policy.

```
student@cp:~/app2$ vim $HOME/app2/allclosed.yaml
```



allclosed.yaml

```
1 ....
2 spec:
```



```

3   podSelector: {}
4   policyTypes:
5   - Ingress
6   # - Egress           #<-- Comment out this line

```

```
student@cp:~/app2$ kubectl replace -f $HOME/app2/allclosed.yaml
```

```
networkpolicy.networking.k8s.io/deny-default replaced
```

6. Test egress access to an outside site. Get the IP address of the **eth0** inside the container while logged in. The IP is 192.168.55.91 in the example below, yours may be different.

```
student@cp:~/app2$ kubectl exec -it -c busy secondapp -- sh
```



On Container

```
/ $ nc -vz www.linux.com 80
```

```
www.linux.com (151.101.185.5:80) open
```

```
/ $ ip a
```

```

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
   inet 127.0.0.1/8 scope host lo
       valid_lft forever preferred_lft forever
   inet6 ::1/128 scope host
       valid_lft forever preferred_lft forever
2: tunl0@NONE: <NOARP> mtu 1480 qdisc noop qlen 1000
   link/ipip 0.0.0.0 brd 0.0.0.0
4: eth0@if59: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue
   link/ether 1e:c8:7d:6a:96:c3 brd ff:ff:ff:ff:ff:ff
   inet 192.168.55.91/32 scope global eth0
       valid_lft forever preferred_lft forever
   inet6 fe80::1cc8:7dff:fe6a:96c3/64 scope link
       valid_lft forever preferred_lft forever

```

```
/ $ exit
```

7. Now add a selector to allow ingress to only the nginx container. Use the IP from the **eth0** range.

```
student@cp:~/app2$ vim $HOME/app2/allclosed.yaml
```



allclosed.yaml

```

1 <output_omitted>
2 policyTypes:
3 - Ingress
4 ingress:           #<-- Add this and following three lines
5 - from:
6   - ipBlock:
7     cidr: 192.168.0.0/16
8 # - Egress

```

8. Recreate the policy, and verify its configuration.

```
student@cp:~/app2$ kubectl replace -f $HOME/app2/allclosed.yaml
```

```
networkpolicy.networking.k8s.io/deny-default replaced
```

```
student@cp:~/app2$ kubectl get networkpolicy
```

NAME	POD-SELECTOR	AGE
deny-default	<none>	3m2s

```
student@cp:~/app2$ kubectl get networkpolicy -o yaml
```

```
apiVersion: v1
items:
- apiVersion: networking.k8s.io/v1
  kind: NetworkPolicy
  metadata:
  <output_omitted>
```

9. Test access to the container both using **curl** as well as **ping**, the IP address to use was found from the IP inside the container. You may need to install **iputils-ping** or other software to use **ping**.

```
student@cp:~/app2$ curl http://192.168.55.91
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

```
student@cp:~/app2$ ping -c5 192.168.55.91
```

```
PING 192.168.55.91 (192.168.55.91) 56(84) bytes of data.
64 bytes from 192.168.55.91: icmp_seq=1 ttl=63 time=1.11 ms
64 bytes from 192.168.55.91: icmp_seq=2 ttl=63 time=0.352 ms
64 bytes from 192.168.55.91: icmp_seq=3 ttl=63 time=0.350 ms
64 bytes from 192.168.55.91: icmp_seq=4 ttl=63 time=0.359 ms
64 bytes from 192.168.55.91: icmp_seq=5 ttl=63 time=0.295 ms

--- 192.168.55.91 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4054ms
rtt min/avg/max/mdev = 0.295/0.495/1.119/0.312 ms
```

10. Update the policy to only allow ingress for TCP traffic on port 80, then test with **curl**, which should work. The ports entry should line up with the from entry a few lines above.

```
student@cp:~/app2$ vim $HOME/app2/allclosed.yaml
```

YA
ML

allclosed.yaml

```
1 <output_omitted>
2   - Ingress
3   ingress:
4   - from:
5     - ipBlock:
6       cidr: 192.168.0.0/16
7     ports:                                #<-- Add this and two following lines
8     - port: 80
```



```
9      protocol: TCP
10 # - Egress
```

```
student@cp:~/app2$ kubectl replace -f $HOME/app2/allclosed.yaml
```

```
networkpolicy.networking.k8s.io/deny-default replaced
```

```
student@cp:~/app2$ curl http://192.168.55.91
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

11. All five pings should fail, with zero received.

```
student@cp:~/app2$ ping -c5 192.168.55.91
```

```
PING 192.168.55.91 (192.168.55.91) 56(84) bytes of data.

--- 192.168.55.91 ping statistics ---
5 packets transmitted, 0 received, 100% packet loss, time 4098ms
```

12. You may want to remove the default-deny policy, in case you want to get to your registry or other pods.

```
student@cp:~/app2$ kubectl delete networkpolicies deny-default
```

```
networkpolicy.networking.k8s.io "deny-default" deleted
```

Exercise 6.6: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

- Understand authentication, authorization, and admission control
- Create & consume Secrets
- Understand ServiceAccounts
- Understand SecurityContexts
- Demonstrate basic understanding of NetworkPolicies

Focus on ensuing you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

1. Create a new deployment which uses the nginx image.
2. Create a new LoadBalancer service to expose the newly created deployment. Test that it works.

3. Create a new NetworkPolicy called netblock which blocks all traffic to pods in this deployment only. Test that all traffic is blocked to deployment.
4. Update the netblock policy to allow traffic to the pod on port 80 only. Test that you can access the default nginx web page.
5. Find and use the `security-review1.yaml` file to create a pod.

```
student@cp:~$ kubectl create -f security-review1.yaml
```

6. View the status of the pod.
7. Use the following commands to figure out why the pod has issues.

```
student@cp:~$ kubectl get pod securityreview
```

```
student@cp:~$ kubectl describe pod securityreview
```

```
student@cp:~$ kubectl logs securityreview
```

8. After finding the errors, log into the container and find the proper id of the nginx user.
9. Edit the yaml and re-create the pod such that the pod runs without error.
10. Create a new serviceAccount called securityaccount.
11. Create a ClusterRole named secrole which only allows create, delete, and list of pods in all apiGroups.
12. Bind the clusterRole to the serviceAccount.
13. Locate the token of the securityaccount. Create a file called `/tmp/securitytoken`. Put only the value of `token:` is equal to, a long string that may start with `eyJh` and be several lines long. Careful that only that string exists in the file.
14. Remove any resources you have added during this review

Chapter 7

Exposing Applications



Exercise 7.1: Exposing Applications: Expose a Service

Overview

In this lab we will explore various ways to expose an application to other pods and outside the cluster. We will add to the NodePort used in previous labs other service options.

1. We will begin by using the default service type ClusterIP. This is a cluster internal IP, only reachable from within the cluster. Begin by viewing the existing services.

```
student@cp:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	NodePort	10.111.26.8	<none>	80:32000/TCP	7h

2. Save then delete the existing service for secondapp. Ensure the same labels, ports, and protocols are used.

```
student@cp:~$ cd $HOME/app2
```

```
student@cp:~/app2$ kubectl get svc secondapp -o yaml > oldservice.yaml
```

```
student@cp:~/app2$ cat oldservice.yaml
```

```
student@cp:~/app2$ kubectl delete svc secondapp
```

```
service "secondapp" deleted
```

3. Recreate the service using a new YAML file. Use the same selector as the previous pod. Examine the new service after creation, note the TYPE and PORT(S).

```
student@cp:~/app2$ vim newservice.yaml
```



newservice.yaml

```

1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: secondapp
5 spec:
6   ports:
7     - port: 80
8     protocol: TCP
9   selector:
10    example: second
11   sessionAffinity: None
12 status:
13   loadBalancer: {}

```

```
student@cp:~/app2$ kubectl create -f newservice.yaml
```

```
service/secondapp created
```

```
student@cp:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	ClusterIP	10.98.148.52	<none>	80/TCP	14s

4. Test access. You should see the default welcome page again.

```
student@cp:~/app2$ curl http://10.98.148.52
```

```

<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```

5. Now create another simple web server. This time use the **httpd** image as the default page is different from **nginx**. Once we are sure it is running we will edit the service selector to point at the new server then back, which could be used as a deployment strategy. Commands have been omitted. You should be able to complete the steps. Refer to previous content otherwise.

```
student@cp:~/app2$ kubectl create deployment newserver --image=httpd
```

6. Locate the newserver labels.
7. Use **kubectl edit** to change the service to use newserver's labels as the selector.
8. Test that the service now shows the new content and not the default **nginx** page.

```
student@cp:~/app2$ curl http://10.98.148.52
```

```
<html><body><h1>It works!</h1></body></html>
```

9. Edit the selector back to the nginx server and test. Then remove the newserver deployment.
10. To expose a port to outside the cluster we will create a NodePort. We had done this in a previous step from the command line. When we create a NodePort it will create a new ClusterIP automatically. Edit the YAML file again. Add type: NodePort. Also add the high-port to match an open port in the firewall as mentioned in the previous chapter. You'll have to delete and re-create as the existing IP is immutable. The NodePort will create a new ClusterIP.

```
student@cp:~/app2$ vim newservice.yaml
```

YAML

newservice.yaml

```
1 ....
2 protocol: TCP
3 nodePort: 32000           #<-- Add this and following line
4 type: NodePort
5 selector:
6   example: second
```

```
student@cp:~/app2$ kubectl delete svc secondapp ; kubectl create -f newservice.yaml
```

```
service "secondapp" deleted
service/secondapp created
```

11. Find the new ClusterIP and ports for the service.

```
student@cp:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	NodePort	10.109.134.221	<none>	80:32000/TCP	4s

12. Test the low port number using the new ClusterIP for the secondapp service.

```
student@cp:~/app2$ curl 10.109.134.221
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

13. Test access from an external node to the host IP and the high container port. Your IP and port will be different. It should work, even with the network policy in place, as the traffic is arriving via a 192.168.0.0 port. If you don't have a terminal on your local system use a browser.

```
student@cp:~/app2$ curl ifconfig.io
```

```
35.184.219.5
```

```
user@laptop:~/Desktop$ curl http://35.184.219.5:32000
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

14. The use of a LoadBalancer makes an asynchronous request to an external provider for a load balancer if one is available. It then creates a NodePort and waits for a response including the external IP. The local NodePort will work even before the load balancer replies. Edit the YAML file and change the type to be LoadBalancer.

```
student@cp:~/app2$ vim newservice.yaml
```



newservice.yaml

```

1 ....
2 - port: 80
3   protocol: TCP
4   nodePort: 32000
5   type: LoadBalancer    #<-- Edit this line
6   selector:
7     example: second

```

```
student@cp:~/app2$ kubectl delete svc secondapp ; kubectl create -f newservice.yaml
```

```

service "secondapp" deleted
service/secondapp created

```

15. As mentioned the cloud provider is not configured to provide a load balancer; the External-IP will remain in pending state. Some issues have been found using this with VirtualBox.

```
student@cp:~/app2$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	8d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	8d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	8d
secondapp	LoadBalancer	10.109.26.21	<pending>	80:32000/TCP	4s

16. Test again local and from a remote node. The IP addresses and ports will be different on your node.

```
serevic@laptop:~/Desktop$ curl http://35.184.219.5:32000
```

```

<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>

```

17. You can also use DNS names provided by **CoreDNS** which dynamically are added when the service is created. Start by logging into the busy container of secondapp.

```
student@cp:~/app2$ kubectl exec -it secondapp -c busy -- sh
```



On Container

- (a) Use the **nslookup** command to find the secondapp service. Then find the registry service we configured to provide container images. If you don't get the expected output try again. About one out of three requests works.

```
/ $ nslookup secondapp
```



```

Server:      10.96.0.10
Address:     10.96.0.10:53

Name:        secondapp.default.svc.cluster.local
Address: 10.96.214.133

*** Can't find secondapp.svc.cluster.local: No answer
*** Can't find secondapp.cluster.local: No answer
*** Can't find secondapp.c.endless-station-188822.internal: No answer
<output_omitted>

```

/ \$ nslookup registry

```

Server:      10.96.0.10
Address:     10.96.0.10:53

Name:        registry.default.svc.cluster.local
Address: 10.110.95.21
<output_omitted>

```

- (b) Lookup the FQDN associated with the DNS server IP displayed by the commands. Your IP may be different.

/ \$ nslookup 10.96.0.10

```

Server:      10.96.0.10
Address:     10.96.0.10:53

10.0.96.10.in-addr.arpa      name = kube-dns.kube-system.svc.cluster.local

```

- (c) Attempt to resolve the service name, which should not bring back any records. Then try with the FQDN. Read through the errors. You'll note that only the default namespaces is checked. You may have to check the FQDN a few times as it doesn't always reply with an answer.

/ \$ nslookup kube-dns

```

Server:      10.96.0.10
Address:     10.96.0.10:53

** server can't find kube-dns.default.svc.cluster.local: NXDOMAIN

*** Can't find kube-dns.svc.cluster.local: No answer
*** Can't find kube-dns.cluster.local: No answer
*** Can't find kube-dns.c.endless-station-188822.internal: No answer

```

/ \$ nslookup kube-dns.kube-system.svc.cluster.local

```

Server:      10.96.0.10
Address:     10.96.0.10:53

Name:        kube-dns.kube-system.svc.cluster.local
Address: 10.96.0.10

*** Can't find kube-dns.kube-system.svc.cluster.local: No answer

```

- (d) Exit out of the container

/ \$ exit

18. Create a new namespace named `multitenant` and a new deployment named `mainapp`. Expose the deployment port 80 using the name `shopping`

```
student@cp:~/app2$ kubectl create ns multitenant
```

```
namespace/multitenant created
```

```
student@cp:~/app2$ kubectl -n multitenant create deployment mainapp --image=nginx
```

```
deployment.apps/mainapp created
```

```
student@cp:~/app2$ kubectl -n multitenant expose deployment mainapp --name=shopping \
--type=NodePort --port=80
```

```
service/shopping exposed
```

19. Log back into the `secondapp` busy container and test access to `mainapp`.

```
student@cp:~/app2$ kubectl exec -it secondapp -c busy -- sh
```



On Container

- (a) Use **nslookup** to determine the address of the new service. Start with using just the service name. Then add the service name and the namespaces. There are a few hiccups, with how busybox and other applications interact with CoreDNS. Your responses may or may not work. Try each a few times.

```
/ $ nslookup shopping
```

```
Server:      10.96.0.10
Address:     10.96.0.10:53

** server can't find shopping.default.svc.cluster.local: NXDOMAIN

*** Can't find shopping.svc.cluster.local: No answer
<output_omitted>
```

```
/ $ nslookup shopping.multitenant
```

```
Server:      10.96.0.10
Address:     10.96.0.10:53

** server can't find shopping.multitenant: NXDOMAIN

*** Can't find shopping.multitenant: No answer
```

```
/ $ nslookup shopping.multitenant.svc.cluster.local
```

```
Server:      10.96.0.10
Address:     10.96.0.10:53

Name:        shopping.multitenant.svc.cluster.local
Address: 10.101.4.142

*** Can't find shopping.multitenant.svc.cluster.local: No answer
```

- (b) Now try to use the service name and then the name with namespace, to see if it works. The DNS using the namespace should work, even if you don't have access to the default page. RBAC could be used to grant access. Check the service ClusterIP returned and it will match the newly created service.



```
/ $ wget shopping
```

```
wget: bad address 'shopping'
```

```
/ $ wget shopping.multitenant
```

```
Connecting to shopping.multitenant (10.101.4.142:80)
wget: can't open 'index.html': Permission denied
```

- (c) As we can see the error is about permissions we will try again, but not try to write locally, but instead to dash (-), which is standard out.

```
~ $ wget -O - shopping.multitenant
```

```
Connecting to shopping.multitenant (10.103.211.64:80)
writing to stdout
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>

....

</html>
-          100% |*****| 612
↪ 0:00:00 ETA
written to stdout
```

Exercise 7.2: Service Mesh and Ingress Controller

If you have a large number of services to expose outside of the cluster, or to expose a low-number port on the host node you can deploy an ingress controller. While nginx and GCE have controllers mentioned a lot in Kubernetes.io, there are many to choose from. Even more functionality and metrics come from the use of a service mesh, such as Istio, Linkerd, Contour, Aspen, or several others.

1. We will install linkerd using their own scripts. There is quite a bit of output. Instead of showing all of it the output has been omitted. Look through the output and ensure that everything gets a green check mark. Some steps may take a few minutes to complete. Each command is listed here to make install easier. As well these steps are in the `setupLinkerd.txt` file.

The most recent versions of linkerd may have some issues. As a result we will use a stable version of Linkerd, such as 2.12.2, or 2.11.4. Feel free to experiment with other versions as time and interest allows. Investigate **linkerd uninstall** and **linkerd uninject** as necessary.

Some commands will take a while to spin up various pods.

```
student@cp:~/app2$ curl -sL run.linkerd.io/install > setup.sh
```

```
student@cp:~/app2$ vim setup.sh
```

```
#!/bin/sh
```

```

set -eu

LINKERD2_VERSION=${LINKERD2_VERSION:-stable-2.12.2} #<-- Edit to earlier stable version
INSTALLROOT=${INSTALLROOT:-"${HOME}/.linkerd2"}

happyexit() {
....

student@cp:~/app2$ sh setup.sh

student@cp:~/app2$ export PATH=$PATH:/home/student/.linkerd2/bin

student@cp:~/app2$ linkerd check --pre

student@cp:~/app2$ linkerd install --crds | kubectl apply -f -

student@cp:~/app2$ linkerd install | kubectl apply -f -

student@cp:~/app2$ linkerd check

student@cp:~/app2$ linkerd viz install | kubectl apply -f -

student@cp:~/app2$ linkerd viz check

student@cp:~/app2$ linkerd viz dashboard &

```

- By default the GUI is on available on the localhost, setup in the final dashboard command. We will need to edit the service and the deployment to allow outside access, in case you are using a cloud provider for the nodes. Edit to remove all characters after equal sign for `-enforced-host`, which is around line 61.

```
student@cp:~/app2$ kubectl -n linkerd-viz edit deploy web
```

YAML

```

1 spec:
2   containers:
3     - args:
4       - -linkerd-controller-api-addr=linkerd-controller-api.linkerd.svc.cluster.local:8085
5       - -linkerd-metrics-api-addr=metrics-api.linkerd-viz.svc.cluster.local:8085
6       - -cluster-domain=cluster.local
7       - -grafana-addr=grafana.linkerd-viz.svc.cluster.local:3000
8       - -controller-namespace=linkerd
9       - -viz-namespace=linkerd-viz
10      - -log-level=info
11      - -enforced-host=                                #<-- Remove everything after equal sign
12      image: cr.l5d.io/linkerd/web:stable-2.12.2
13      imagePullPolicy: IfNotPresent

```

- Now edit the http nodePort and type to be a NodePort.

```
student@cp:~/app2$ kubectl edit svc web -n linkerd-viz
```

YAML

```

1 ....
2 ports:
3   - name: http
4     nodePort: 31500                                #<-- Add line with an easy to remember port
5     port: 8084
6   ....
7   sessionAffinity: None
8   type: NodePort                                    #<-- Edit type to be NodePort

```


YAML

```

9 status:
10   loadBalancer: {}
11   ...

```

4. Test access using a local browser to your public IP. Your IP will be different than the one shown below.

```
student@cp:~/app2$ curl ifconfig.io
```

```
104.197.159.20
```

5. From your local system open a browser and go to the public IP and the high-number nodePort, such as 31500. It will not be the localhost port seen in output.




Figure 7.1: Main Linkerd Page

6. In order for linkerd to pay attention to an object we need to add an annotation. The **linkerd inject** command will do this for us. Generate YAML and pipe it to **linkerd** then pipe again to **kubectl**. Expect an error about how the object was created, but the process will work. The command can run on one line if you omit the back-slash.

```
student@cp:~/app2$ kubectl -n multitenant get deploy mainapp -o yaml | \
  linkerd inject - | kubectl apply -f -
```

```
<output_omitted>
```

7. Check the GUI, you should see that the **multitenant** namespaces and pods are now meshed, and the name is a link.
8. Generate some traffic to the pods, and watch the traffic via the GUI.

```
student@cp:~$ kubectl -n multitenant get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
shopping	NodePort	10.102.8.205	<none>	80:32518/TCP	4h54m

```
student@cp:~$ curl 10.102.8.205
```

```

<!DOCTYPE html>
<html>
<head>

```

```
<title>Welcome to nginx!</title>
<output_omitted>
```

Namespaces		HTTP Metrics				
Control Plane		Namespace ↑	↑ Meshed	↑ Success Rate	↑ RPS	↑ P50 Latency
DEFAULT		default	0/9	---	---	---
OADS		kube-node-lease	0/0	---	---	---
Cron Jobs		kube-public	0/0	---	---	---
Daemon Sets		kube-system	0/11	---	---	---
Deployments		linkerd	5/5	100.00% ●	3.07	1 ms
Jobs		linkerd-viz	6/6	100.00% ●	8.2	5 ms
Pods		multitenant	1/1	100.00% ●	0.3	1 ms
Replica Sets						
Replication Controllers						
Stateful Sets						

Figure 7.2: Now shows meshed

9. Scale up the mainapp deployment. Generate traffic to get metrics for all the pods.

```
student@cp:~$ kubectl -n multitenant scale deploy mainapp --replicas=5
```

```
deployment.apps/mainapp scaled
```

```
student@cp:~$ curl 10.102.8.205 #Several times
```

10. Explore some of the other information provided by the GUI.

Jobs	linkerd	5/5	100.00% ●	3	1 ms
Pods	linkerd-viz	6/6	100.00% ●	4.58	7 ms
Replica Sets	multitenant	5/5	100.00% ●	0.45	1 ms
Replication Controllers					

Figure 7.3: Five meshed pods

11. Linkerd does not come with an ingress controller, so we will add one to help manage traffic. We will leverage a **Helm** chart to install an ingress controller. Search the hub to find that there are many available.

```
student@cp:~$ helm search hub ingress
```

```
URL                                CHART VERSION
APP VERSION                        DESCRIPTION
https://artifacthub.io/packages/helm/k8s-as-hel... 1.0.2
v1.0.0                             Helm Chart representing a single Ingress Kubern...
https://artifacthub.io/packages/helm/openstack-... 0.2.1
v0.32.0                             OpenStack-Helm Ingress Controller
<output_omitted>
```

```
https://artifacthub.io/packages/helm/api/ingres... 3.29.1
0.45.0 Ingress controller for Kubernetes using NGINX a...
https://artifacthub.io/packages/helm/wener/ingr... 3.31.0
0.46.0 Ingress controller for Kubernetes using NGINX a...
https://artifacthub.io/packages/helm/nginx/ngin... 0.9.2
1.11.2 NGINX Ingress Controller
<output_omitted>
```

12. We will use a popular ingress controller provided by **NGINX**.

```
student@cp:~$ helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx
```

```
"ingress-nginx" has been added to your repositories
```

```
student@cp:~$ helm repo update
```

```
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "ingress-nginx" chart repository
Update Complete. -Happy Helming!-
```

13. Download and edit the `values.yaml` file and change it to use a DaemonSet instead of a Deployment. This way there will be a pod on every node to handle traffic if using an external load balancer.

```
student@cp:~$ helm fetch ingress-nginx/ingress-nginx --untar
```

```
student@cp:~$ cd ingress-nginx
```

```
student@cp:~/ingress-nginx$ ls
```

```
CHANGELOG.md Chart.yaml OWNERS README.md ci templates values.yaml
```

```
student@cp:~/ingress-nginx$ vim values.yaml
```

YA
ML

values.yaml

```
1 ....
2 ## DaemonSet or Deployment
3 ##
4 kind: DaemonSet                                #<-- Change to DaemonSet, around line 150
5
6 ## Annotations to be added to the controller Deployment or DaemonSet
7 ....
```

14. Now install the controller using the chart. Note the use of the dot (.) to look in the current directory.

```
student@cp:~/ingress-nginx$ helm install myingress .
```

```
NAME: myingress
LAST DEPLOYED: Wed May 19 22:24:27 2021
NAMESPACE: default
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The ingress-nginx controller has been installed.
It may take a few minutes for the LoadBalancer IP to be available.
You can watch the status by running
'kubectl --namespace default get services -o wide -w myingress-ingress-nginx-controller'
```

An example Ingress that makes use of the controller:
<output_omitted>

15. We now have an ingress controller running, but no rules yet. View the resources that exist. Use the **-w** option to watch the ingress controller service show up. After it is available use **ctrl-c** to quit and move to the next command.

```
student@cp:~$ kubectl get ingress --all-namespaces
```

No resources found

```
student@cp:~$ kubectl --namespace default get services -o wide -w myingress-ingress-nginx-controller
```

NAME	PORT(S)	AGE	TYPE	SELECTOR	CLUSTER-IP	EXTERNAL-IP
myingress-ingress-nginx-controller	80:32558/TCP,443:30219/TCP	47s	LoadBalancer	app.kubernetes.io/component=controller, app.kubernetes.io/instance=myingress, app.kubernetes.io/name=ingress-nginx	10.104.227.79	<pending>

```
student@cp:~$ kubectl get pod --all-namespaces |grep nginx
```

default	myingress-ingress-nginx-controller-mrqt5	1/1	Running	0	20s
default	myingress-ingress-nginx-controller-pkdxm	1/1	Running	0	62s
default	nginx-b68dd9f75-h6ww7	1/1	Running	0	21h

16. Now we can add rules which match HTTP headers to services. Remember to check the course files from the tarball.

```
student@cp:~$ vim ingress.yaml
```

YAML

ingress.yaml

```
1 apiVersion: networking.k8s.io/v1
2 kind: Ingress
3 metadata:
4   name: ingress-test
5   annotations:
6     kubernetes.io/ingress.class: "nginx"
7   namespace: default
8 spec:
9   rules:
10  - host: www.example.com
11    http:
12      paths:
13      - backend:
14          service:
15            name: secondapp
16            port:
17              number: 80
18          path: /
19          pathType: ImplementationSpecific
20 status:
21   loadBalancer: {}
```

17. Create then verify the ingress is working. If you don't pass a matching header you should get a 404 error.

```
student@cp:~$ kubectl create -f ingress.yaml
```

ingress.networking.k8s.io/ingress-test created

```
student@cp:~$ kubectl get ingress
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
ingress-test	<none>	www.example.com		80	5s

```
student@cp:~$ kubectl get pod -o wide |grep myingress
```

myingress-ingress-nginx-controller-mrqt5	1/1	Running	0	8m9s	192.168.219.118
cp <none>	<none>				
myingress-ingress-nginx-controller-pkdxm	1/1	Running	0	8m9s	192.168.219.118
cp <none>	<none>				

```
student@cp:~/ingress-nginx$ curl 192.168.219.118
```

```
<html>
<head><title>404 Not Found</title></head>
<body>
<center><h1>404 Not Found</h1></center>
<hr><center>nginx</center>
</body>
</html>
```

18. Check the ingress service and expect another 404 error, don't use the admission controller.

```
student@cp:~/ingress-nginx$ kubectl get svc |grep ingress
```

myingress-ingress-nginx-controller	LoadBalancer	10.104.227.79	<pending>
80:32558/TCP,443:30219/TCP	10m		
myingress-ingress-nginx-controller-admission	ClusterIP	10.97.132.127	<none>
443/TCP	10m		

```
student@cp:~/ingress-nginx$ curl 10.104.227.79
```

```
<html>
<head><title>404 Not Found</title></head>
<body>
<center><h1>404 Not Found</h1></center>
<hr><center>nginx</center>
</body>
</html>
```

19. Now pass a matching header and you should see the default web server page.

```
student@cp:~/ingress-nginx$ curl -H "Host: www.example.com" http://10.104.227.79
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
<output_omitted>
```

20. We can add an annotation to the ingress pods for Linkerd. You will get some warnings, but the command will work.

```
student@cp:~/ingress-nginx$ kubectl get ds myingress-ingress-nginx-controller -o yaml | \
linkerd inject --ingress - | kubectl apply -f -
```

```
daemonset "myingress-ingress-nginx-controller" injected
Warning: resource daemonsets/myingress-ingress-nginx-controller is missing the
```

```
kubectl.kubernetes.io/last-applied-configuration annotation which is required
by kubectl apply. kubectl apply should only be used on resources created
declaratively by either kubectl create --save-config or kubectl apply. The
missing annotation will be patched automatically.
daemonset.apps/myingress-ingress-nginx-controller configured
```

21. Go to the Top page, change the namespace to default and the resource to daemonset/myingress-ingress-nginx-controller. Press start then pass more traffic to the ingress controller and view traffic metrics via the GUI. Let top run so we can see another page added in an upcoming step.




Figure 7.4: Ingress Traffic

22. At this point we would keep adding more and more web servers. We'll configure one more, which would then could be a process continued as many times as desired. Begin by deploying another **nginx** server. Give it a label and expose port 80.

```
student@cp:~/app2$ kubectl create deployment thirdpage --image=nginx
```

```
deployment.apps "thirdpage" created
```

23. Assign a label for the ingress controller to match against. Your pod name is unique, you can use the **Tab** key to complete the name.

```
student@cp:~/app2$ kubectl label pod thirdpage-<tab> example=third
```

24. Expose the new server as a NodePort.

```
student@cp:~/app2$ kubectl expose deployment thirdpage --port=80 --type=NodePort
```

```
service/thirdpage exposed
```

25. Now we will customize the installation. Run a bash shell inside the new pod. Your pod name will end differently. Install **vim** or an editor inside the container then edit the `index.html` file of nginx so that the title of the web page will be Third Page. Much of the command output is not shown below.

```
student@cp:~/app2$ kubectl exec -it thirdpage-<Tab> -- /bin/bash
```



On Container

```
root@thirdpage-:/# apt-get update

root@thirdpage-:/# apt-get install vim -y

root@thirdpage-:/# vim /usr/share/nginx/html/index.html

<!DOCTYPE html>
<html>
<head>
<title>Third Page</title>      #<-- Edit this line
<style>
<output_omitted>

root@thirdpage-:/# exit
```

Edit the ingress rules to point the thirdpage service. It may be easiest to copy the existing host stanza and edit the host and name.

26. `student@cp:~/app2$ kubectl edit ingress ingress-test`



ingress-test

```
1 ....
2 spec:
3   rules:
4     - host: thirdpage.org
5       http:
6         paths:
7           - backend:
8               service:
9                 name: thirdpage
10                port:
11                  number: 80
12              path: /
13              pathType: ImplementationSpecific
14     - host: www.example.com
15       http:
16         paths:
17           - backend:
18               service:
19                 name: secondapp
20                port:
21                  number: 80
22              path: /
23              pathType: ImplementationSpecific
24   status:
25   ....
```

27. Test the second Host: setting using **curl** locally as well as from a remote system, be sure the <title> shows the non-default page. Use the main IP of either node. The Linkerd GUI should show a new T0 line, if you select the small blue box with an arrow you will see the traffic is going to thirdpage.

`student@cp:~/app2$ curl -H "Host: thirdpage.org" http://10.128.0.7/`

```
<!DOCTYPE html>
<html>
<head>
```

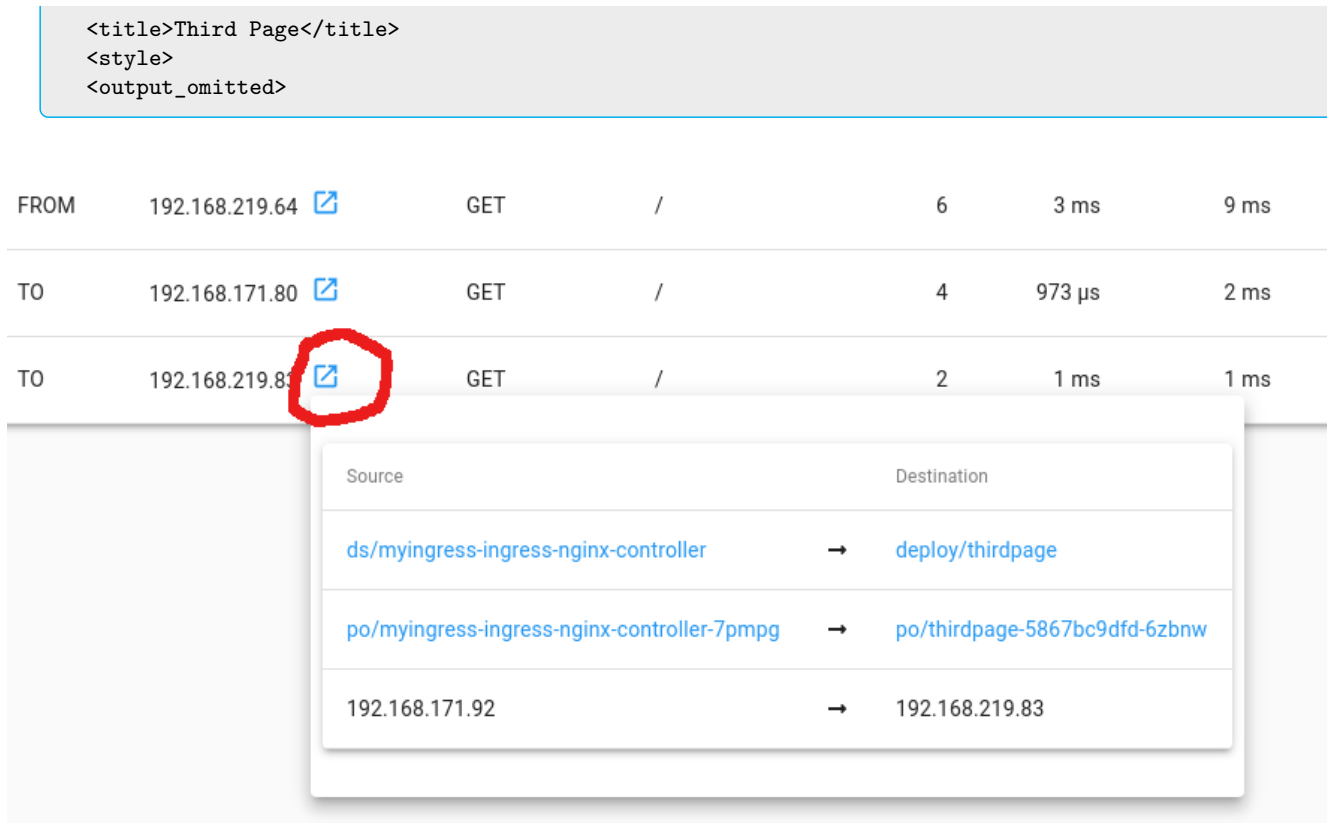


Figure 7.5: Linkerd Top Metrics

28. Consider how you would edit the ingress rules to point at a different server, as we did when editing the service selector. Deploy the **httpd** server again and test changing traffic from `thirdpage` over to a new `httpd` sever using only ingress rule edits.

✍ Exercise 7.3: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

- Use Kubernetes primitives to implement common deployment strategies (e.g. blue/gren or canary)
- Use the Helm package manager to deploy existing packages
- Provide and troubleshoot access to applications via services
- Use Ingress rules to expose applications

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

Using the browser and the three URL locations allowed by the exam, find and bookmark working YAML examples to do the following:

1. Create a new pod called `webone`, running the `nginx` service. Expose port 80.

2. Create a new service named `webone-svc`. The service should be accessible from outside the cluster.
3. Update both the pod and the service with selectors so that traffic for to the service IP shows the web server content.
4. Change the type of the service such that it is only accessible from within the cluster. Test that exterior access no longer works, but access from within the node works.
5. Deploy another pod, called `webtwo`, this time running the `wlniao/website` image. Create another service, called `webtwo-svc` such that only requests from within the cluster work. Note the default page for each server is distinct.
6. Install and configure an ingress controller such that requests for `webone.com` see the `nginx` default page, and requests for `webtwo.org` see the `wlniao/website` default page.
7. Remove any resources created in this review.

Chapter 8

Application Troubleshooting



Exercise 8.1: Troubleshooting: Monitor Applications

Overview

Troubleshooting can be difficult in a multi-node, decoupled and transient environment. Add in the rapid pace of change and it becomes more difficult. Instead of focusing and remembering a particular error and the fix it may be more useful to learn a flow of troubleshooting and revisit assumptions until the pace of change slows and various areas further mature.

1. View the `secondapp` pod, it should show as `Running`. This may not mean the application within is working properly, but that the pod is running. The restarts are due to the command we have written to run. The pod exists when done, and the controller restarts another container inside. The count depends on how long the labs have been running.

```
student@cp/app2:~$ cd
student@cp:~$ kubectl get pods secondapp
```

NAME	READY	STATUS	RESTARTS	AGE
secondapp	2/2	Running	49	2d

2. Look closer at the pod. Working slowly through the output check each line. If you have issues, are other pods having issues on the same node or volume? Check the state of each container. Both `busy` and `webserver` should report as `Running`. Note `webserver` has a restart count of zero while `busy` has a restart count of 49. We expect this as, in our case, the pod has been running for 49 hours.

```
student@cp:~$ kubectl describe pod secondapp
```

```
Name:          secondapp
Namespace:     default
Node:          worker-wdrq/10.128.0.2
Start Time:    Fri, 13 Apr 2018 20:34:56 +0000
Labels:        example=second
Annotations:    <none>
Status:        Running
IP:            192.168.55.91
Containers:
  webserver:
    <output_omitted>
    State:      Running
```

```

    Started:      Fri, 13 Apr 2018 20:34:58 +0000
    Ready:        True
    Restart Count: 0
<output_omitted>

    busy:
<output_omitted>

    State:        Running
    Started:      Sun, 15 Apr 2018 21:36:20 +0000
    Last State:    Terminated
    Reason:        Completed
    Exit Code:     0
    Started:      Sun, 15 Apr 2018 20:36:18 +0000
    Finished:      Sun, 15 Apr 2018 21:36:18 +0000
    Ready:        True
    Restart Count: 49
    Environment:   <none>

```

3. There are three values for conditions. Check that the pod reports Initialized, Ready and scheduled.

```

<output_omitted>
Conditions:
  Type           Status
  Initialized     True
  Ready           True
  PodScheduled    True
<output_omitted>

```

4. Check if there are any events with errors or warnings which may indicate what is causing any problems.

```

Events:
  Type    Reason      Age           From              Message
  ----    -
  Normal  Pulling     34m (x50 over 2d) kubelet, worker-wdrq pulling
image "busybox"
  Normal  Pulled      34m (x50 over 2d) kubelet, worker-wdrq Successfully
pulled image "busybox"
  Normal  Created     34m (x50 over 2d) kubelet, worker-wdrq Created
container
  Normal  Started     34m (x50 over 2d) kubelet, worker-wdrq Started
container

```

5. View each container log. You may have to sift errors from expected output. Some containers may have no output at all, as is found with busy.

```
student@cp:~$ kubectl logs secondapp webserver
```

```

192.168.55.0 - - [13/Apr/2018:21:18:13 +0000] "GET / HTTP/1.1" 200
612 "-" "curl/7.47.0" "-"
192.168.55.0 - - [13/Apr/2018:21:20:35 +0000] "GET / HTTP/1.1" 200
612 "-" "curl/7.53.1" "-"
127.0.0.1 - - [13/Apr/2018:21:25:29 +0000] "GET" 400 174 "-" "-" "-"
127.0.0.1 - - [13/Apr/2018:21:26:19 +0000] "GET index.html" 400 174
"-" "-" "-"
<output_omitted>

```

```
student@cp:~$ kubectl logs secondapp busy
```

```
student@cp:~$
```

6. Check to make sure the container is able to use DNS and communicate with the outside world. Remember we still have limited the UID for `secondapp` to be UID **2000**, which may prevent some commands from running. It can also prevent an application from completing expected tasks, and other errors.

```
student@cp:~$ kubectl exec -it secondapp -c busy -- sh
```



On Container

```
/ $ nslookup www.linuxfoundation.org
```

```
/ $ nslookup www.linuxfoundation.org
Server:          10.96.0.10
Address:         10.96.0.10:53

Non-authoritative answer:
Name:   www.linuxfoundation.org
Address: 23.185.0.2

*** Can't find www.linuxfoundation.org: No answer
```

```
/ $ cat /etc/resolv.conf
```

```
nameserver 10.96.0.10
search default.svc.cluster.local svc.cluster.local
cluster.local c.endless-station-188822.internal
google.internal
options ndots:5
```

7. Test access to a remote node using **nc (NetCat)**. There are several options to **nc** which can help troubleshoot if the problem is the local node, something between nodes or in the target. In the example below the connect never completes and a **control-c** was used to interrupt.

```
/ $ nc www.linux.com 25
```

```
^Cpunt!
```

8. Test using an IP address in order to narrow the issue to name resolution. In this case the IP in use is a well known IP for Google's DNS servers. The following example shows that Internet name resolution is working, but our UID issue prevents access to the `index.html` file.

```
/ $ wget http://www.linux.com/
```

```
Connecting to www.linux.com (151.101.45.5:80)
Connecting to www.linux.com (151.101.45.5:443)
wget: can't open 'index.html': Permission denied

/ $ exit
```

9. Make sure traffic is being sent to the correct Pod. Check the details of both the service and endpoint. Pay close attention to ports in use as a simple typo can prevent traffic from reaching the proper pod. Make sure labels and selectors don't have any typos as well.

```
student@cp:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	10d
nginx	ClusterIP	10.108.95.67	<none>	443/TCP	10d
registry	ClusterIP	10.105.119.236	<none>	5000/TCP	10d
secondapp	LoadBalancer	10.109.26.21	<pending>	80:32000/TCP	1d

```
thirdpage    NodePort    10.109.250.78    <none>    80:31230/TCP    1h
```

```
student@cp:~$ kubectl get svc secondapp -o yaml
```

```
<output_omitted>
clusterIP: 10.109.26.21
externalTrafficPolicy: Cluster
ports:
- nodePort: 32000
  port: 80
  protocol: TCP
  targetPort: 80
selector:
  example: second
<output_omitted>
```

10. Verify an endpoint for the service exists and has expected values, including namespaces, ports and protocols.

```
student@cp:~$ kubectl get ep
```

```
NAME           ENDPOINTS          AGE
kubernetes     10.128.0.3:6443    10d
nginx          192.168.55.68:443  10d
registry       192.168.55.69:5000 10d
secondapp      192.168.55.91:80   1d
thirdpage      192.168.241.57:80  1h
```

```
student@cp:~$ kubectl get ep secondapp -o yaml
```

```
apiVersion: v1
kind: Endpoints
metadata:
  creationTimestamp: 2018-04-14T05:37:32Z
<output_omitted>
```

11. If the containers, services and endpoints are working the issue may be with an infrastructure service like **kube-proxy**. Ensure it's running, then look for errors in the logs. As we have two nodes we will have two proxies to look at. As we built our cluster with **kubeadm** the proxy runs as a container. On other systems you may need to use **journalctl** or look under `/var/log/kube-proxy.log`.

```
student@cp:~$ ps -elf |grep kube-proxy
```

```
<output_omitted>
4 S root      2652  2621  0  80   0 - 186908 -      15:23 ?          00:00:03
↪ /usr/local/bin/kube-proxy
--config=/var/lib/kube-proxy/config.conf --hostname-override=cp
0 S student   31656  9707  0  80   0 - 3715 pipe_w 19:02 pts/0    00:00:00 grep --color=auto
↪ kube-proxy
```

```
student@cp:~$ journalctl -a | grep proxy
```

```
Apr 15 15:44:43 worker-nzjr audit[742]: AVC apparmor="STATUS"
operation="profile_load" profile="unconfined" \
name="/usr/lib/lxd/lxd-bridge-proxy" pid=742 comm="apparmor_parser"
Apr 15 15:44:43 worker-nzjr kernel: audit: type=1400
audit(1523807083.011:11): apparmor="STATUS" \
operation="profile_load" profile="unconfined" \
name="/usr/lib/lxd/lxd-bridge-proxy" pid=742 comm="apparmor_parser"
Apr 15 15:45:17 worker-nzjr kubelet[1248]: I0415 15:45:17.153670
1248 reconciler.go:217] operationExecutor.VerifyControllerAttachedVolume\
started for volume "xtables-lock" \
```

```
(UniqueName: "kubernetes.io/host-path/e701fc01-38f3-11e8-a142-\
42010a800003-xtables-lock") \
pod "kube-proxy-t8k4w" (UID: "e701fc01-38f3-11e8-a142-42010a800003")
```

12. Look at both of the proxy logs. Lines which begin with the character **I** are info, **E** are errors. In this example the last message says access to listing an endpoint was denied by RBAC. It was because a default installation via Helm wasn't RBAC aware. This is only an example, you (hopefully) won't see any errors on your lab nodes.

If not using command line completion, view the possible pod names first.

```
student@cp:~$ kubectl -n kube-system get pod
```

```
student@cp:~$ kubectl -n kube-system logs kube-proxy-fsdf
```

```
I0405 17:28:37.091224      1 feature_gate.go:190] feature gates: map[]
W0405 17:28:37.100565      1 server_others.go:289] Flag proxy-mode=""
unknown, assuming iptables proxy
I0405 17:28:37.101846      1 server_others.go:138] Using iptables Proxier.
I0405 17:28:37.121601      1 server_others.go:171] Tearing down
inactive rules.
<output_omitted>
E0415 15:45:17.086081      1 reflector.go:205] \
k8s.io/kubernetes/pkg/client/informers/informers_generated/
internalversion/factory.go:85: \
Failed to list *core.Endpoints: endpoints is forbidden: \
User "system:serviceaccount:kube-system:kube-proxy" cannot \
list endpoints at the cluster scope:\
[clusterrole.rbac.authorization.k8s.io "system:node-proxier" not found, \
clusterrole.rbac.authorization.k8s.io "system:basic-user" not found, \
clusterrole.rbac.authorization.k8s.io \
"system:discovery" not found]
```

13. Check that the proxy is creating the expected rules for the problem service. Find the destination port being used for the service, **32000** in this case.

```
student@cp:~$ sudo iptables-save |grep secondapp
```

```
-A KUBE-NODEPORTS -p tcp -m comment --comment "default/secondapp:" \
-m tcp --dport 32000 -j KUBE-MARK-MASQ
-A KUBE-NODEPORTS -p tcp -m comment --comment "default/secondapp:" \
-m tcp --dport 32000 -j KUBE-SVC-DAASHM5XQZF5XI3E
-A KUBE-SERVICES ! -s 192.168.0.0/16 -d 10.109.26.21/32 -p tcp \
-m comment --comment "default/secondapp: \
cluster IP" -m tcp --dport 80 -j KUBE-MARK-MASQ
-A KUBE-SERVICES -d 10.109.26.21/32 -p tcp -m comment --comment \
"default/secondapp: cluster IP" -m tcp \
--dport 80 -j KUBE-SVC-DAASHM5XQZF5XI3E
<output_omitted>
```

14. Ensure the proxy is working by checking the port targeted by **iptables**. If it fails open a second terminal and view the proxy logs when making a request as it happens.

```
student@cp:~$ curl localhost:32000
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

✍ Exercise 8.2: Update YAML File

The API continues to change, both as objects mature as well as new features and settings being added. As a result YAML that may have worked on a previous version of Kubernetes may need to be updated due to API deprecations.

1. Find the `brokendeploy.yaml` file from the course tarball. Try to create the deployment, which worked a few releases of Kubernetes prior. Create a new deployment from the command line for an understanding of what the current object requires. Edit the broken deployment YAML until you can deploy it and the pod runs.

```
student@cp:~$ find $HOME -name brokendeploy.yaml

student@cp:~$ cp <path_from_above> .

student@cp:~$ kubectl create -f brokendeploy.yaml
```

✍ Exercise 8.3: OPTIONAL LAB: Conformance Testing

The **cncf.io** group is in the process of formalizing what is considered to be a conforming Kubernetes cluster. While that project matures there is an existing tool provided by **Sonobuoy** <https://sonobuoy.io/> which can be useful. We will need to make sure a newer version of **Golang** is installed for it to work. You can download the code from github and look around with git or with go, depending on which tool you are most familiar. **Things change quickly these steps may not work....today**

1. Download a compiled binary. A shorter URL is shown first, then the longer, just in case the link changes and you need to navigate. They should download the same file.

```
student@cp:~$ curl -sL0 https://tinyurl.com/4cdeyz55

student@cp:~$ mv 4cdeyz55 sonobuoy.tar.gz

student@cp:~$ tar -xvf sonobuoy.tar.gz
```

```
LICENSE
sonobuoy
```

```
student@cp:~$ curl -sL0 \
https://github.com/vmware-tanzu/sonobuoy/releases/download/v0.56.8/sonobuoy_0.56.8_linux_amd64.tar.gz
```

2. Run the test. Use the `--wait` option, which will capture the screen until the test finishes. This could take a while to finish, such as an hour. You should get some output indicating testing objects being created.

```
student@cp:~$ sudo mv sonobuoy /usr/local/bin/

student@cp:~$ sonobuoy run --wait
```

```
INFO[0000] create request issued          name=sonobuoy namespace=
resource=namespaces
INFO[0000] create request issued          name=sonobuoy-serviceaccount
namespace=sonobuoy resource=serviceaccounts
INFO[0000] create request issued          name=sonobuoy-serviceaccount
-sonobuoy namespace= resource=clusterrolebindings
INFO[0000] create request issued          name=sonobuoy-serviceaccount
-sonobuoy namespace= resource=clusterroles
INFO[0000] create request issued          name=sonobuoy-config-cm
namespace=sonobuoy resource=configmaps
INFO[0000] create request issued          name=sonobuoy-plugins-cm
namespace=sonobuoy resource=configmaps
INFO[0000] create request issued          name=sonobuoy namespace=
sonobuoy resource=pods
INFO[0000] create request issued          name=sonobuoy-aggregator
```



```

namespace=sonobuoy resource=services
14:45:09          PLUGIN      NODE      STATUS      RESULT      PROGRESS
14:45:09          e2e        global    running      Passed: 0, Failed: 0, Remaining:354
14:45:09    systemd-logs      cp-1    complete
14:45:09    systemd-logs    worker-1    complete
14:45:09
14:45:09 Sonobuoy is still running. Runs can take 60 minutes or more depending
on cluster and plugin configuration.

<output_omitted>

16:38:29          e2e        global    running      Passed:345, Failed: 3, Remaining: 6
16:38:49          e2e        global    complete     Passed:351, Failed: 3, Remaining: 0
16:38:49 Sonobuoy plugins have completed. Preparing results for download.
16:39:09          e2e        global    complete    failed    Passed:351, Failed: 3, Remaining: 0
16:39:09    systemd-logs      cp-1    complete    passed
16:39:09    systemd-logs    worker-1    complete    passed
16:39:09 Sonobuoy has completed. Use `sonobuoy retrieve` to get results.

```

3. If you don't want to wait for the full results, open a second terminal session to the cp and use the **status** and **logs** sub commands.

```
student@cp:~$ sonobuoy status
```

```

          PLUGIN  STATUS      RESULT  COUNT
PROGRESS
          e2e    running      1    Passed: 31, Failed: 0, Remaining:323
systemd-logs  complete      2

Sonobuoy is still running. Runs can take 60 minutes or more depending on
cluster and plugin configuration.

```

```
student@cp:~$ sonobuoy logs
```

```

namespace="sonobuoy" pod="sonobuoy-systemd-logs-daemon-set-34bef90c2026439b-x7q7g"
container="sonobuoy-worker"
time="2022-07-23T14:44:51Z" level=trace msg="Invoked command single-node with args
[] and flags [level=trace logtostderr=true sleep=-1 v=6]"
time="2022-07-23T14:44:51Z" level=info msg="Waiting for waitfile"
waitfile=/tmp/sonobuoy/results/done
time="2022-07-23T14:44:51Z" level=info msg="Starting to listen on port 8099 for
progress updates and will relay them to
https://[192.168.193.162]:8080/api/v1/progress/by-node/cp-1/systemd-logs"
....

```

4. Wait until the results are ready, or log into the container and look for results.

```
student@cp:~$ sonobuoy retrieve
```

```
202207231444_sonobuoy_7d9c5bb5-e547-4827-bb2e-07e12f8fed25.tar.gz
```

5. View the results

```
student@cp:~$ sonobuoy results 202207231444_sonobuoy_7d9c5bb5-e547-4827-bb2e-07e12f8fed25.tar.gz
```

```

Plugin: e2e
Status: failed
Total: 6971
Passed: 351
Failed: 3
Skipped: 6617

```

```
Failed tests:
[sig-scheduling] SchedulerPredicates [Serial] validates that NodeSelector is respected if not
↳ matching [Conformance]
<output_omitted>
```

6. Delete the pods and namespaces created by Sonobuoy.

```
student@cp:~$ sonobuoy delete --wait
```

✍ Exercise 8.4: Domain Review



Very Important

The source pages and content in this review could change at any time. **IT IS YOUR RESPONSIBILITY TO CHECK THE CURRENT INFORMATION.**

Revisit the CKAD domain list on [Curriculum Overview](#) and locate some of the topics we have covered in this chapter.

- Understand API deprecations
- Use provided tools to monitor Kubernetes applications
- Debugging in Kubernetes

Focus on ensuring you have all the necessary files and processes understood first. Repeat the review until you are sure you have bookmarks of necessary YAML samples and can complete each step quickly, and ensure each object is running properly.

1. Find and use the `troubleshoot-review1.yaml` file to create a deployment. The **create** command will fail. Edit the file to fix issues such that a single pod runs for at least a minute without issue. There are several things to fix.

```
student@cp:~$ kubectl create -f troubleshoot-review1.yaml
```

```
<Fix any errors found here>
```

When fixed it should look like this:

```
student@cp:~$ kubectl get deploy igottrouble
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
igottrouble	1/1	1	1	5m13s

2. Remove any resources created during this review.